

Introduction

Trusted Computing Base (TCB)

- ▶ Base components that enforce the fundamental protection mechanisms on a computing system
 - ◆ Hardware
 - ◆ Firmware
 - ◆ Software
- ▶ TCB vulnerabilities potentially affect the security of the entire system

TCB by TCSEC (Trusted Computer System Evaluation Criteria, aka Orange Book)

The totality of protection mechanisms within a computing system – including hardware, firmware, and software – the combination of which is responsible for enforcing a computer security policy.

A TCB consists of one or more components that together enforce a unified security policy over a product or system.

The ability of a trusted computing base to correctly enforce a security policy depends solely on the mechanisms within the TCB and on the correct input by system administrative personnel of parameters (e.g., a user's clearance) related to the security policy.

TCB by MITRE

Nibaldi, G. H. *Specification of a trusted computing base (TCB)*. MITRE CORP
BEDFORD MA, 1979.

A TCB is a hardware and software access control mechanism that establishes a protection environment to control the sharing of information in computer systems. A TCB is an implementation of a reference monitor, [...], that controls when and how data is accessed.

TCB fundamental components

- ▷ CPU security mechanisms
 - ◆ Protection rings
 - ◆ Virtualization
 - ◆ Other mechanisms
 - E.g. Intel SGX enclaves, etc.
- ▷ Operating system security model
 - ◆ Computational model
 - ◆ Access rights and privileges

TEE (Trusted Execution Environment)

- ▷ Isolated, secure execution environment
- ▷ CPU support
 - ♦ ARM TrustZone
- ▷ TEE implementations
 - ♦ On-board Credentials (Microsoft/Nokia)
 - ♦ <t-base (Trustonic)
 - ♦ SecuriTEE (Solacia)
 - ♦ QSEE (Qualcomm's Secure Execution Environment)
 - ♦ SierraTEE (Sierrawave, open-source)
 - ♦ OP-TEE (Linaro, open-source)

Can you trust the operating system?

- ▶ Can you trust your operating system if you do not control (or trust) the way it booted?
- ▶ Secure bootstrapping
 - ◆ TPM attestation
 - ◆ UEFI secure boot
- ▶ Remote attestation
 - ◆ TPM attestation

Can you trust the operating system?

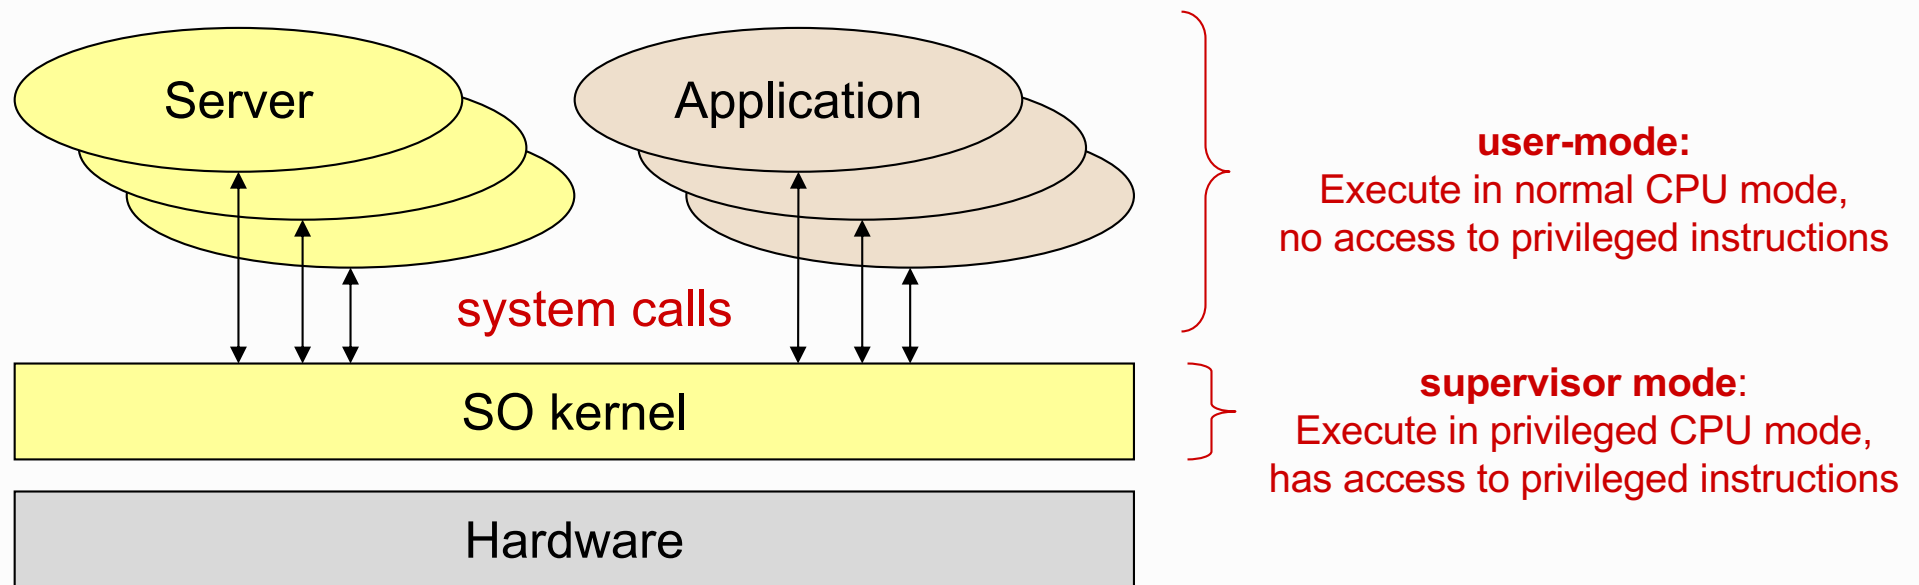
- ▷ How can you protect your computation if you don't trust the operating system?
- ▷ Intel SGX (Secure Guard eXtensions)
 - ♦ Allow user applications to protect code and data from others within enclaves
 - ♦ Enclaves are not observable by code running with different privileges
 - OS kernels, hypervisors, etc.

Protection from untrusted code: sandboxes

- ▷ Executing applications have a set of privileges and a view over a set of resources
- ▷ Sandboxes allow the execution of applications with less privileges or less resources
 - ◆ e.g. forbid remote communications
 - ◆ e.g. hide the majority of the file system
 - ◆ e.g. allow volatile system changes

Security in Operating Systems

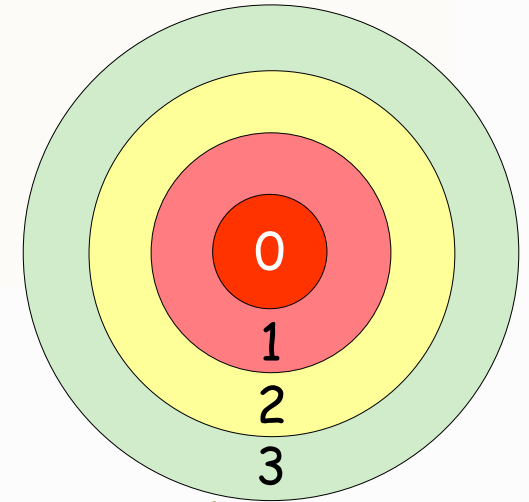
Operating system



▷ Kernel mission

- ♦ Virtualize the hardware
 - Computational model
- ♦ Enforce protection policies and provide protection mechanisms
 - Against involuntary mistakes
 - Against non-authorized activities

Protection rings



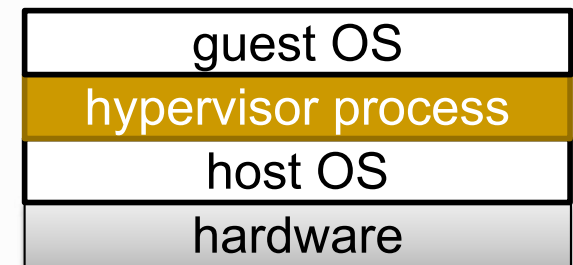
- ▷ Different levels of privilege
 - ◆ Forming a set of concentric rings
 - ◆ Used by CPU's to prevent non-privileged code from running privileged instructions
 - e.g. IN/OUT, TLB manipulation

- ▷ Nowadays processors have 4 rings
 - ◆ But OS's usually use only two of them
 - 0 (supervisor/kernel mode) and 3 (user-mode)

- ▷ Transfer of control between rings requires special gates
 - ◆ The ones that are used by system calls (syscalls)

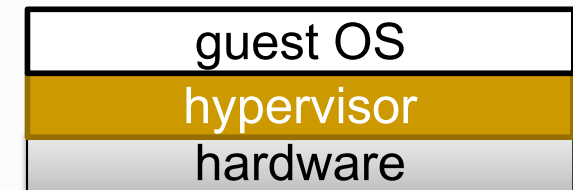
Virtual machines and hypervisors

- ▶ Emulation of a particular (virtual) hardware with the existing one (real)



- ▶ Hosted virtualization

- ◆ The hypervisor is a process of a given OS (host)
- ◆ The VM runs inside the virtualizer (guest OS)



- ▶ Bare-metal virtualization

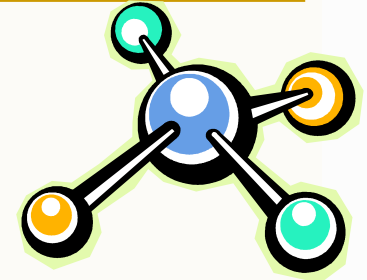
- ◆ The hypervisor runs on top of the host hardware

Execution of virtual machines

- ▷ Common approach for hosted virtualization
 - ◆ Software-based virtualization
 - ◆ Direct execution of guest user-mode code
 - ◆ Binary, on-the-fly translation of privileged code (full virtualization)
 - Guest OS kernels remain unchanged
 - No direct access to the host hardware

- ▷ Hardware-assisted virtualization (bare-metal)
 - ◆ Full virtualization
 - ◆ There is a ring -1 below ring 0
 - Hypervisor (or Virtual Machine Monitor, VMM)
 - ◆ It can virtualize hardware for many ring 0 kernels
 - No need of binary translation
 - Guest OS's run faster

Computational model



- ▷ Set of entities (objects) managed by the OS kernel
 - ♦ High-level abstractions supported transparently by low-level mechanisms
- ▷ Processes
- ▷ User identifiers
 - ♦ Users
 - ♦ Groups
- ▷ Virtual memory
- ▷ Files and file systems
 - ♦ Directories
 - ♦ Files
 - ♦ Special files
- ▷ Communication channels
 - ♦ Pipes
 - ♦ Sockets
 - ♦ Etc.
- ▷ Physical devices
 - ♦ Storage
 - Tapes
 - Magnetic disks
 - Optical disks
 - SSD
 - ♦ Network interfaces
 - Wired, wireless
 - ♦ Human-computer interfaces
 - Keyboards
 - Graphical screens
 - Text consoles
 - Mice
 - ♦ Serial/parallel I/O interfaces
 - USB
 - Serial & parallel ports
 - Bluetooth

Computational model: User identifiers



- ▷ For the OS kernel a user is a number
 - ◆ Established during a login operation
 - ◆ User ID (UID)

- ▷ All activities are executed on a computer on behalf of a UID
 - ◆ The UID allows the kernel to assert what is allowed/denied to processes
 - ◆ Linux: UID 0 is omnipotent (root)
 - Administration activities are usually executed with UID 0
 - ◆ Windows: concept of privileges
 - For administration, system configuration, etc.
 - There is no unique, well-known identifier for an administrator
 - Administration privileges can be bound to several UIDs
 - Usually through administration groups
 - Administrators, Power Users, Backup Operators
 - ◆ Linux: concept of capabilities (similar to privileges)

Computational model:

Group identifiers



- ▷ Groups also have an identifier
 - ◆ A group is a set of users
 - ◆ A group can be defined by including other groups
 - ◆ Group ID (GID)
- ▷ A user can belong to several groups
 - ◆ Actual user rights = UID rights + rights of his groups' GIDs
- ▷ In Linux all activities are executed on behalf of a set of groups
 - ◆ Primary group
 - Typically used for setting file protection
 - ◆ Secondary groups

Computational model:

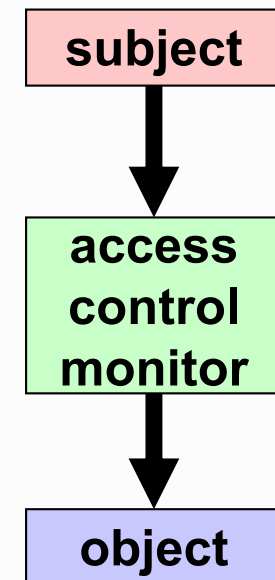
Processes

- ▷ A process defines the context of an activity
 - ♦ For taking security-related decisions
 - ♦ For other purposes (e.g. scheduling)

- ▷ Security-related context
 - ♦ Identity (UID and GIDs)
 - Fundamental for enforcing access control
 - ♦ Resources being used
 - Open files
 - Including communication channels
 - Reserved virtual memory areas
 - CPU time used

Access control

- ▷ The OS kernel is an access control monitor
 - ◆ Controls all interactions of subjects with protected objects
- ▷ Objects
 - ◆ Hardware
 - ◆ Entities of the computational model
- ▷ Subjects
 - ◆ Usually local processes
 - Through the system call API
 - A system call (or syscall) is not an ordinary function call
 - ◆ But also messages from other hosts



Mandatory access controls

- ▷ OS kernels have plenty mandatory access control policies
 - ♦ They are part of the computational model logic
 - ♦ They cannot be overruled not even by administrators
 - Unless they change the OS kernel behavior

- ▷ Examples:
 - ♦ Kernel runs in CPU privileged modes, user applications run in non-privileged modes
 - ♦ Separation of virtual memory areas
 - ♦ Inter-process signaling
 - ♦ Interpretation of files' access control protections

Protection with ACLs (Access Control Lists)

- ▷ Each object has an ACL
 - ♦ It says which subjects can do what
- ▷ An ACL can be discretionary or mandatory
 - ♦ When mandatory it cannot be modified
 - ♦ When discretionary it can be tailored
- ▷ An ACL is checked when an activity, on behalf of a subject, wants to manipulate the object
 - ♦ If the manipulation request is not authorized by the ACL, the access is denied
 - ♦ The OS kernel is responsible for enforcing ACL-based protection

Protection with capabilities

- ▷ Less common in normal OS kernels
 - ♦ Though there are some good examples
- ▷ Example: open file descriptors
 - ♦ Applications' processes indirectly manipulate (open) files through file descriptors kept by the OS kernel
 - File descriptors are referenced using integer indexes (aka file descriptors for simplicity...)
 - The OS kernel has full control over the contents of open file descriptors
 - ♦ Access to open file descriptors can only be granted to other processes through the OS kernel
 - Not really a usual operation, but possible!
 - ♦ Changes in the protection of files does not impact existing open file descriptors
 - The access rights are evaluated and memorized when the file is open

Unix file protection ACLs:

Fixed-structure, discretionary ACL

- ▷ Each file system object has an ACL
 - ◆ Binding 3 rights to 3 subjects
 - ◆ Only the owner can update the ACL
- ▷ Rights: **R W X**
 - ◆ Read (file data) / List directory
 - ◆ Write (file data) / create or remove files or subdirectories
 - ◆ Execute / use as process' current working directory
- ▷ Subjects:
 - ◆ An UID (owner)
 - ◆ A GID
 - ◆ Others

Windows NTFS file protection:

Variable-size, discretionary ACLs

- ▷ Each file system object has an ACL and a owner
 - ♦ 13 types of access rights
 - ♦ Variable-size list of subjects
 - ♦ Owner can be an UID or a GID
 - ♦ Owner has no special rights over the object or its ACL
 - But usually file creators are their initial owners and have Change Permissions rights

- ▷ Subjects:
 - ♦ Users (UIDs)
 - ♦ Groups (GIDs)
 - The group "Everyone" stands for anybody

- ▷ Access rights:

File	Directory (folder)
Read (data)	List (files / folders)
Write (data)	Create (files)
Append (data)	Create (folders)
Execute	Traverse
Delete (file)	Delete (folder)
	Delete (files and subfolders)
Read attributes / extended attributes	
Write attributes / extended attributes	
Read permissions	
Change permissions	
Take ownership	

Unix file protection ACLs:

Special protection bits

▷ Set-UID bit

```
creator:Pictures$ ls -la /usr/bin/passwd  
-rwsr-xr-x 1 root root 59640 Mar 22 2019 /usr/bin/passwd
```

- ♦ Is used to change the UID of processes executing the file

▷ Set-GID bit

```
creator:Pictures$ ls -la /usr/bin/at  
-rwsr-sr-x 1 daemon daemon 51464 Feb 20 2018 /usr/bin/at
```

- ♦ Is used to change the GID of processes executing the file

▷ Sticky bit

```
creator:Pictures$ ls -la /tmp  
total 108  
drwxrwxrwt 25 root root 4096 Dec 15 13:12 .
```

- ♦ Hint to keep the file/directory as much as possible in memory cache

Privilege elevation:

Set-UID mechanism

- ▷ It is used to change the UID of a process running a program stored on a Set-UID file
 - ♦ If a program file is owned by UID **X** and the set-UID bit of its ACL is set, then it will be executed in a process with UID **X**
 - Independently of the UID of the subject that executed the program
- ▷ Used to allow normal users to execute privileged tasks encapsulated in administration programs
 - ♦ Change the user's password (**passwd**)
 - ♦ Change to super-user mode (**su**, **sudo**)
 - ♦ Mount devices (**mount**)

Privilege elevation:

Set-UID mechanism (cont.)

- ▷ Effective UID / Real UID
 - ♦ **Real UID** is the UID of the process creator
 - App launcher
 - ♦ **Effective UID** is the UID of the process
 - The one that really matters for defining the rights of the process

- ▷ UID change
 - ♦ **Ordinary application**
 - eUID = rUID = UID of process that executed **exec**
 - eUID cannot be changed (unless = 0)
 - ♦ **Set-UID application**
 - eUID = UID of **exec**'d application file, rUID = initial process UID
 - eUID can revert to rUID
 - ♦ **rUID cannot change**

Privilege elevation: Set-UID/Set-GID decision flowchart

▷ exec (path, ...)

- ♦ File referred by path has Set-UID?
- ♦ Yes
 - ID = path owner
 - Change the process effective UID to ID
- ♦ No
 - Do nothing
- ♦ File referred by path has Set-GID?
- ♦ Yes
 - ID = path GID
 - Change the process GID to ID only
- ♦ No
 - Do nothing

Privilege elevation: sudo mechanism

- ▷ Administration by root is not advised
 - ♦ One "identity", many people
 - ♦ Who did what?
- ▷ Preferable approach
 - ♦ Administration role (uid = 0), many users assume it
 - Sudoers
 - Defined by a configuration file used by sudo
- ▷ **sudo** is a Set-UID application with UID = 0
 - ♦ Logging can take place on each command ran with sudo

Privilege reduction: chroot mechanism (or jail)

- ▷ Used to reduce the visibility of a file system
 - ◆ Each process descriptor has a **root i-node number**
 - From which absolute pathname resolution takes place
 - ◆ **chroot** changes it to an arbitrary directory
 - The process' file system view gets reduced

- ▷ Used to protect the file system from potentially problematic applications
 - ◆ e.g. public servers, downloaded applications
 - ◆ But it is not bullet proof!

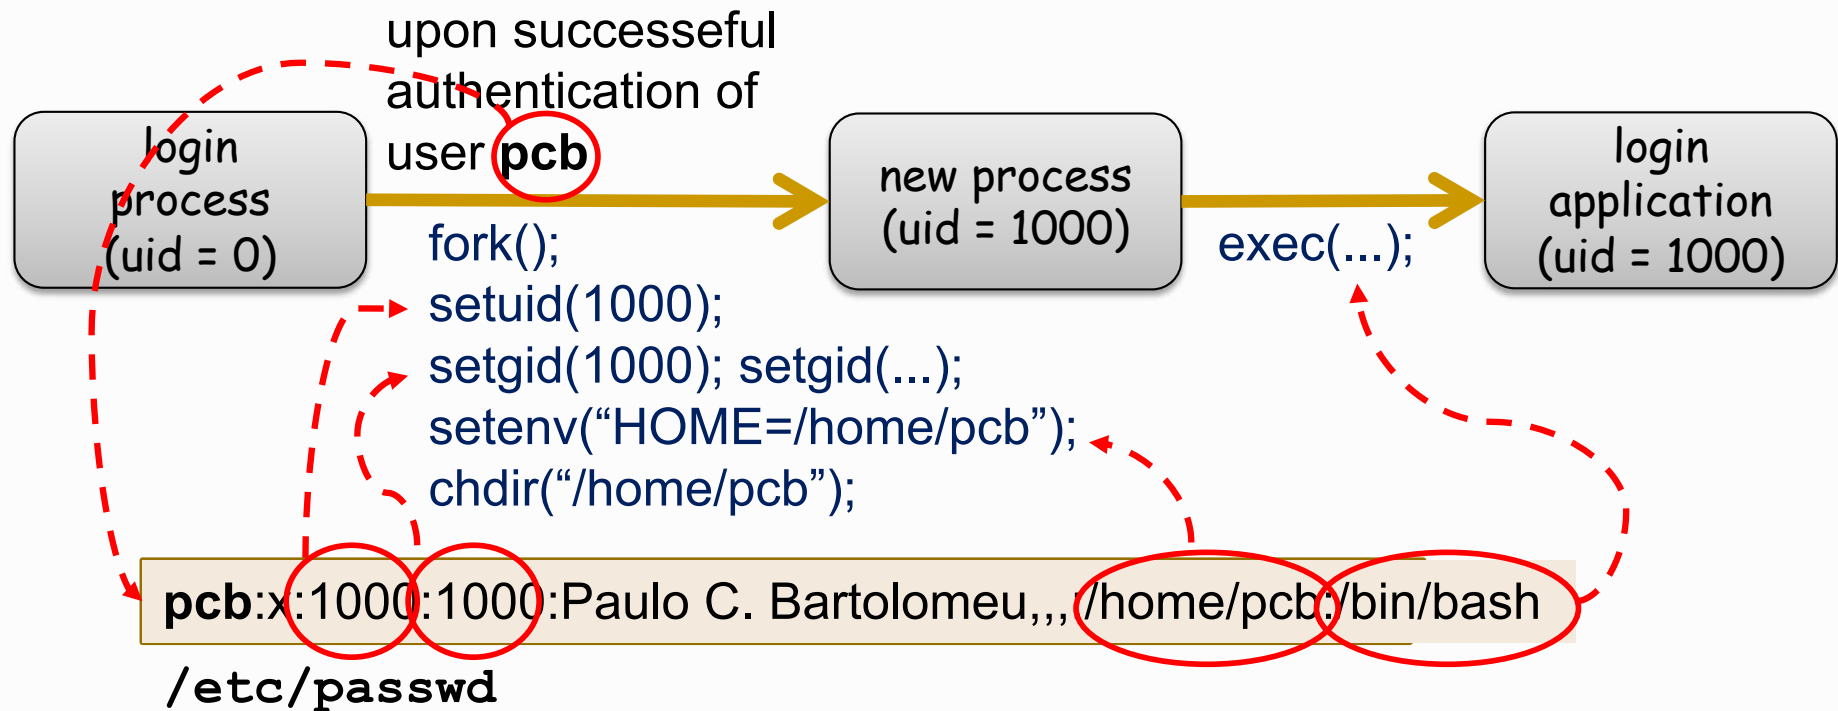
Linux login:

Not an OS kernel operation

- ▷ A privileged login application presents an interface for getting users' credentials
 - ♦ A username/password pair
 - ♦ Biometric data
 - ♦ Smartcard and activation PIN
- ▷ The login application validates the credentials and fetches the appropriate UID and GIDs for the user
 - ♦ And starts an initial user application on a process with those identifiers
 - In a Linux console this application is a shell (sh, bash, csh, tcsh, zsh, etc.)
 - ♦ When this process ends the login application reappears
- ▷ Thereafter all processes created by the user have its identifiers
 - ♦ Inherited through forks

Linux: from login to session processes

- ▶ The login process must be a privileged process
 - ◆ Has to create processes with arbitrary UID and GIDs
 - The ones of the entity logging in



Login in Linux:

Password validation process

- ▷ Username is used to fetch a UID/GID pair from `/etc/passwd`
 - ◆ And a set of additional GIDs in the `/etc/group` file
- ▷ Supplied password is transformed using a digest function
 - ◆ Currently configurable, for creating a new user (`/etc/login.defs`)
 - ◆ Its identification is stored along with the transformed password
- ▷ The result is checked against a value stored in `/etc/shadow`
 - ◆ Indexed again by the username
 - ◆ If they match, the user was correctly authenticated
- ▷ File protections
 - ◆ `/etc/passwd` and `/etc/group` can be read by anyone
 - This is fundamental, for instance, for listing directories (why?)
 - ◆ `/etc/shadow` can only be read by root
 - Protection against dictionary attacks