

Introduction

Trusted Computing Base (TCB)

- ▶ Base components that enforce the fundamental protection mechanisms on a computing system
 - ◆ Hardware
 - ◆ Firmware
 - ◆ Software
- ▶ TCB vulnerabilities potentially affect the security of the entire system

TCB by TCSEC (Trusted Computer System Evaluation Criteria, aka Orange Book)

The totality of protection mechanisms within a computing system – including hardware, firmware, and software – the combination of which is responsible for enforcing a computer security policy.

A TCB consists of one or more components that together enforce a unified security policy over a product or system.

The ability of a trusted computing base to correctly enforce a security policy depends solely on the mechanisms within the TCB and on the correct input by system administrative personnel of parameters (e.g., a user's clearance) related to the security policy.

TCB by MITRE

Nibaldi, G. H. *Specification of a trusted computing base (TCB)*. MITRE CORP
BEDFORD MA, 1979.

A TCB is a hardware and software access control mechanism that establishes a protection environment to control the sharing of information in computer systems. A TCB is an implementation of a reference monitor, [...], that controls when and how data is accessed.

TCB fundamental components

- ▷ CPU security mechanisms
 - ◆ Protection rings
 - ◆ Virtualization
 - ◆ Other mechanisms
 - E.g. Intel SGX enclaves, etc.
- ▷ Operating system security model
 - ◆ Computational model
 - ◆ Access rights and privileges

TEE (Trusted Execution Environment)

- ▷ Isolated, secure execution environment
- ▷ CPU support
 - ♦ ARM TrustZone
- ▷ TEE implementations
 - ♦ On-board Credentials (Microsoft/Nokia)
 - ♦ <t-base (Trustonic)
 - ♦ SecuriTEE (Solacia)
 - ♦ QSEE (Qualcomm's Secure Execution Environment)
 - ♦ SierraTEE (Sierrawave, open-source)
 - ♦ OP-TEE (Linaro, open-source)

Can you trust the operating system?

- ▶ Can you trust your operating system if you do not control (or trust) the way it booted?
- ▶ Secure bootstrapping
 - ◆ TPM attestation
 - ◆ UEFI secure boot
- ▶ Remote attestation
 - ◆ TPM attestation

Can you trust the operating system?

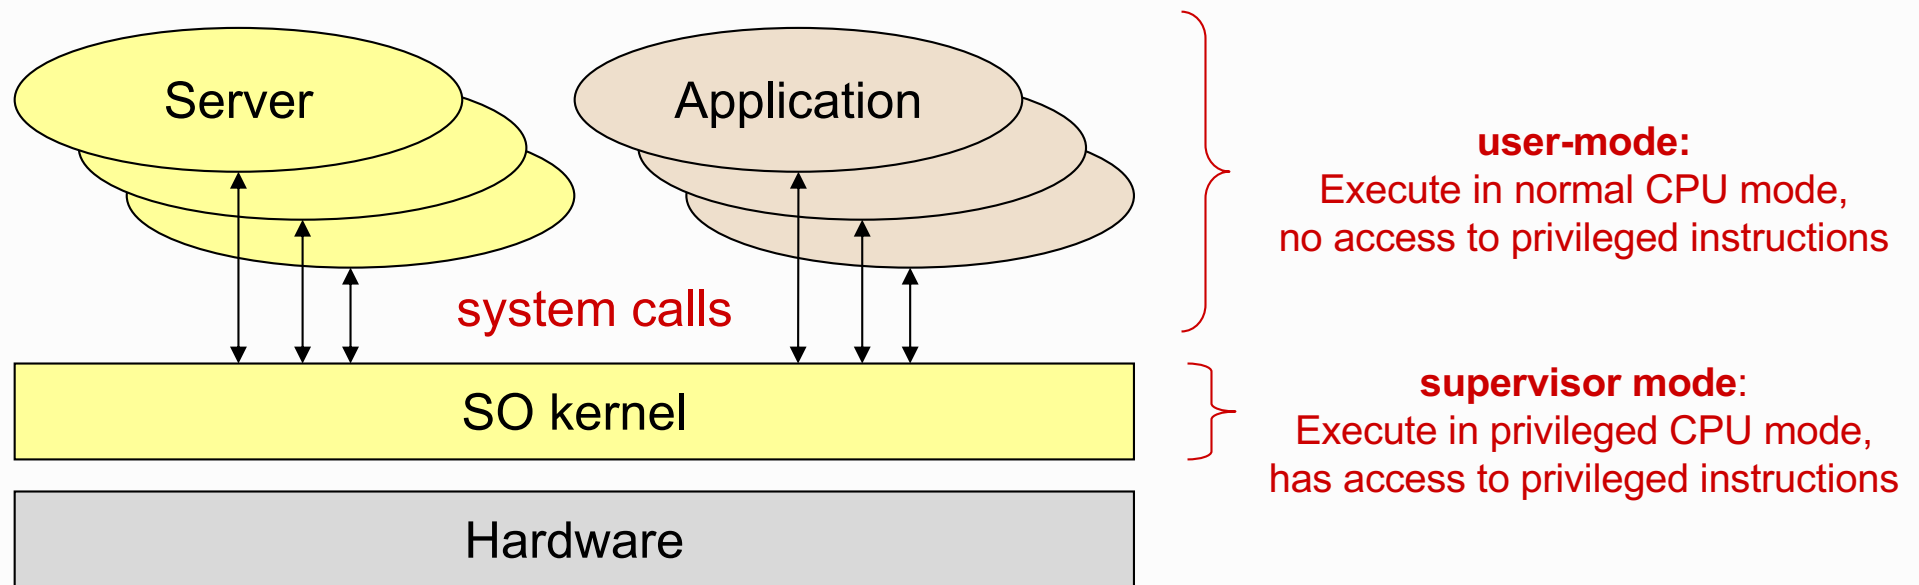
- ▷ How can you protect your computation if you don't trust the operating system?
- ▷ Intel SGX (Secure Guard eXtensions)
 - ♦ Allow user applications to protect code and data from others within enclaves
 - ♦ Enclaves are not observable by code running with different privileges
 - OS kernels, hypervisors, etc.

Protection from untrusted code: sandboxes

- ▷ Executing applications have a set of privileges and a view over a set of resources
- ▷ Sandboxes allow the execution of applications with less privileges or less resources
 - ◆ e.g. forbid remote communications
 - ◆ e.g. hide the majority of the file system
 - ◆ e.g. allow volatile system changes

Security in Operating Systems

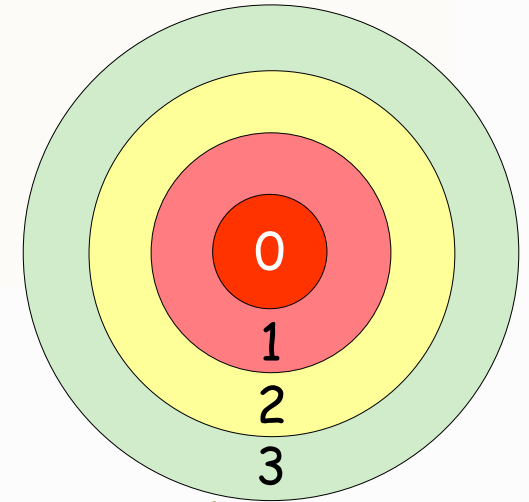
Operating system



► Kernel mission

- ♦ Virtualize the hardware
 - Computational model
- ♦ Enforce protection policies and provide protection mechanisms
 - Against involuntary mistakes
 - Against non-authorized activities

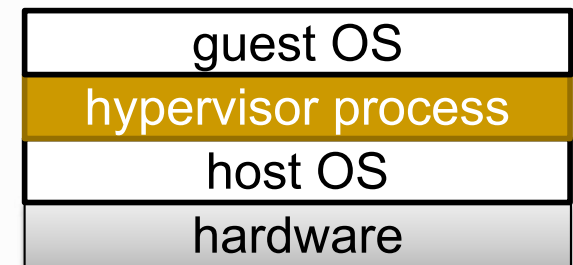
Protection rings



- ▷ Different levels of privilege
 - ◆ Forming a set of concentric rings
 - ◆ Used by CPU's to prevent non-privileged code from running privileged instructions
 - e.g. IN/OUT, TLB manipulation
- ▷ Nowadays processors have 4 rings
 - ◆ But OS's usually use only two of them
 - 0 (supervisor/kernel mode) and 3 (user-mode)
- ▷ Transfer of control between rings requires special gates
 - ◆ The ones that are used by system calls (syscalls)

Virtual machines and hypervisors

- ▶ Emulation of a particular (virtual) hardware with the existing one (real)



- ▶ Hosted virtualization

- ◆ The hypervisor is a process of a given OS (host)
- ◆ The VM runs inside the virtualizer (guest OS)



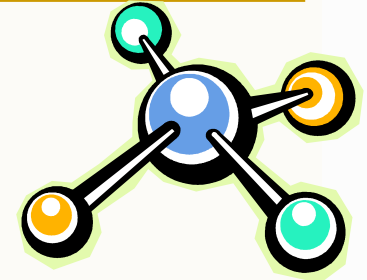
- ▶ Bare-metal virtualization

- ◆ The hypervisor runs on top of the host hardware

Execution of virtual machines

- ▷ Common approach for hosted virtualization
 - ◆ Software-based virtualization
 - ◆ Direct execution of guest user-mode code
 - ◆ Binary, on-the-fly translation of privileged code (full virtualization)
 - Guest OS kernels remain unchanged
 - No direct access to the host hardware
- ▷ Hardware-assisted virtualization (bare-metal)
 - ◆ Full virtualization
 - ◆ There is a ring -1 below ring 0
 - Hypervisor (or Virtual Machine Monitor, VMM)
 - ◆ It can virtualize hardware for many ring 0 kernels
 - No need of binary translation
 - Guest OS's run faster

Computational model



- ▷ Set of entities (objects) managed by the OS kernel
 - ♦ High-level abstractions supported transparently by low-level mechanisms
- ▷ Processes
- ▷ User identifiers
 - ♦ Users
 - ♦ Groups
- ▷ Virtual memory
- ▷ Files and file systems
 - ♦ Directories
 - ♦ Files
 - ♦ Special files
- ▷ Communication channels
 - ♦ Pipes
 - ♦ Sockets
 - ♦ Etc.
- ▷ Physical devices
 - ♦ Storage
 - Tapes
 - Magnetic disks
 - Optical disks
 - SSD
 - ♦ Network interfaces
 - Wired, wireless
 - ♦ Human-computer interfaces
 - Keyboards
 - Graphical screens
 - Text consoles
 - Mice
 - ♦ Serial/parallel I/O interfaces
 - USB
 - Serial & parallel ports
 - Bluetooth

Computational model: User identifiers



- ▷ For the OS kernel a user is a number
 - ◆ Established during a login operation
 - ◆ User ID (UID)

- ▷ All activities are executed on a computer on behalf of a UID
 - ◆ The UID allows the kernel to assert what is allowed/denied to processes
 - ◆ Linux: UID 0 is omnipotent (root)
 - Administration activities are usually executed with UID 0
 - ◆ Windows: concept of privileges
 - For administration, system configuration, etc.
 - There is no unique, well-known identifier for an administrator
 - Administration privileges can be bound to several UIDs
 - Usually through administration groups
 - Administrators, Power Users, Backup Operators
 - ◆ Linux: concept of capabilities (similar to privileges)

Computational model:

Group identifiers



- ▷ Groups also have an identifier
 - ♦ A group is a set of users
 - ♦ A group can be defined by including other groups
 - ♦ Group ID (GID)
- ▷ A user can belong to several groups
 - ♦ Actual user rights = UID rights + rights of his groups' GIDs
- ▷ In Linux all activities are executed on behalf of a set of groups
 - ♦ Primary group
 - Typically used for setting file protection
 - ♦ Secondary groups

Computational model:

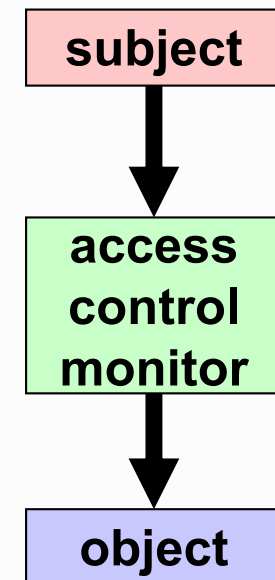
Processes

- ▷ A process defines the context of an activity
 - ♦ For taking security-related decisions
 - ♦ For other purposes (e.g. scheduling)

- ▷ Security-related context
 - ♦ Identity (UID and GIDs)
 - Fundamental for enforcing access control
 - ♦ Resources being used
 - Open files
 - Including communication channels
 - Reserved virtual memory areas
 - CPU time used

Access control

- ▷ The OS kernel is an access control monitor
 - ◆ Controls all interactions of subjects with protected objects
- ▷ Objects
 - ◆ Hardware
 - ◆ Entities of the computational model
- ▷ Subjects
 - ◆ Usually local processes
 - Through the system call API
 - A system call (or syscall) is not an ordinary function call
 - ◆ But also messages from other hosts



Mandatory access controls

- ▷ OS kernels have plenty mandatory access control policies
 - ♦ They are part of the computational model logic
 - ♦ They cannot be overruled not even by administrators
 - Unless they change the OS kernel behavior

- ▷ Examples:
 - ♦ Kernel runs in CPU privileged modes, user applications run in non-privileged modes
 - ♦ Separation of virtual memory areas
 - ♦ Inter-process signaling
 - ♦ Interpretation of files' access control protections

Protection with ACLs (Access Control Lists)

- ▷ Each object has an ACL
 - ♦ It says which subjects can do what
- ▷ An ACL can be discretionary or mandatory
 - ♦ When mandatory it cannot be modified
 - ♦ When discretionary it can be tailored
- ▷ An ACL is checked when an activity, on behalf of a subject, wants to manipulate the object
 - ♦ If the manipulation request is not authorized by the ACL, the access is denied
 - ♦ The OS kernel is responsible for enforcing ACL-based protection

Protection with capabilities

- ▷ Less common in normal OS kernels
 - ♦ Though there are some good examples
- ▷ Example: open file descriptors
 - ♦ Applications' processes indirectly manipulate (open) files through file descriptors kept by the OS kernel
 - File descriptors are referenced using integer indexes (aka file descriptors for simplicity...)
 - The OS kernel has full control over the contents of open file descriptors
 - ♦ Access to open file descriptors can only be granted to other processes through the OS kernel
 - Not really a usual operation, but possible!
 - ♦ Changes in the protection of files does not impact existing open file descriptors
 - The access rights are evaluated and memorized when the file is open

Unix file protection ACLs:

Fixed-structure, discretionary ACL

- ▷ Each file system object has an ACL
 - ◆ Binding 3 rights to 3 subjects
 - ◆ Only the owner can update the ACL
- ▷ Rights: **R W X**
 - ◆ Read (file data) / List directory
 - ◆ Write (file data) / create or remove files or subdirectories
 - ◆ Execute / use as process' current working directory
- ▷ Subjects:
 - ◆ An UID (owner)
 - ◆ A GID
 - ◆ Others

Windows NTFS file protection:

Variable-size, discretionary ACLs

- ▷ Each file system object has an ACL and a owner
 - ♦ 13 types of access rights
 - ♦ Variable-size list of subjects
 - ♦ Owner can be an UID or a GID
 - ♦ Owner has no special rights over the object or its ACL
 - But usually file creators are their initial owners and have Change Permissions rights

- ▷ Subjects:
 - ♦ Users (UIDs)
 - ♦ Groups (GIDs)
 - The group "Everyone" stands for anybody

- ▷ Access rights:

File	Directory (folder)
Read (data)	List (files / folders)
Write (data)	Create (files)
Append (data)	Create (folders)
Execute	Traverse
Delete (file)	Delete (folder)
	Delete (files and subfolders)
Read attributes / extended attributes	
Write attributes / extended attributes	
Read permissions	
Change permissions	
Take ownership	

Unix file protection ACLs:

Special protection bits

▷ Set-UID bit

```
creator:Pictures$ ls -la /usr/bin/passwd  
-rwsr-xr-x 1 root root 59640 Mar 22 2019 /usr/bin/passwd
```

- ♦ Is used to change the UID of processes executing the file

▷ Set-GID bit

```
creator:Pictures$ ls -la /usr/bin/at  
-rwsr-sr-x 1 daemon daemon 51464 Feb 20 2018 /usr/bin/at
```

- ♦ Is used to change the GID of processes executing the file

▷ Sticky bit

```
creator:Pictures$ ls -la /tmp  
total 108  
drwxrwxrwt 25 root root 4096 Dec 15 13:12 .
```

- ♦ Hint to keep the file/directory as much as possible in memory cache

Privilege elevation:

Set-UID mechanism

- ▷ It is used to change the UID of a process running a program stored on a Set-UID file
 - ♦ If a program file is owned by UID **X** and the set-UID bit of its ACL is set, then it will be executed in a process with UID **X**
 - Independently of the UID of the subject that executed the program
- ▷ Used to allow normal users to execute privileged tasks encapsulated in administration programs
 - ♦ Change the user's password (**passwd**)
 - ♦ Change to super-user mode (**su**, **sudo**)
 - ♦ Mount devices (**mount**)

Privilege elevation:

Set-UID mechanism (cont.)

- ▷ Effective UID / Real UID
 - ♦ **Real UID** is the UID of the process creator
 - App launcher
 - ♦ **Effective UID** is the UID of the process
 - The one that really matters for defining the rights of the process
- ▷ UID change
 - ♦ **Ordinary application**
 - eUID = rUID = UID of process that executed **exec**
 - eUID cannot be changed (unless = 0)
 - ♦ **Set-UID application**
 - eUID = UID of **exec**'d application file, rUID = initial process UID
 - eUID can revert to rUID
 - ♦ **rUID cannot change**

Privilege elevation: Set-UID/Set-GID decision flowchart

▷ exec (path, ...)

- ♦ File referred by path has Set-UID?
- ♦ Yes
 - ID = path owner
 - Change the process effective UID to ID
- ♦ No
 - Do nothing
- ♦ File referred by path has Set-GID?
- ♦ Yes
 - ID = path GID
 - Change the process GID to ID only
- ♦ No
 - Do nothing

Privilege elevation: sudo mechanism

- ▷ Administration by root is not advised
 - ♦ One "identity", many people
 - ♦ Who did what?
- ▷ Preferable approach
 - ♦ Administration role (uid = 0), many users assume it
 - Sudoers
 - Defined by a configuration file used by sudo
- ▷ **sudo** is a Set-UID application with UID = 0
 - ♦ Logging can take place on each command ran with sudo

Privilege reduction: chroot mechanism (or jail)

- ▷ Used to reduce the visibility of a file system
 - ◆ Each process descriptor has a **root i-node number**
 - From which absolute pathname resolution takes place
 - ◆ **chroot** changes it to an arbitrary directory
 - The process' file system view gets reduced

- ▷ Used to protect the file system from potentially problematic applications
 - ◆ e.g. public servers, downloaded applications
 - ◆ But it is not bullet proof!

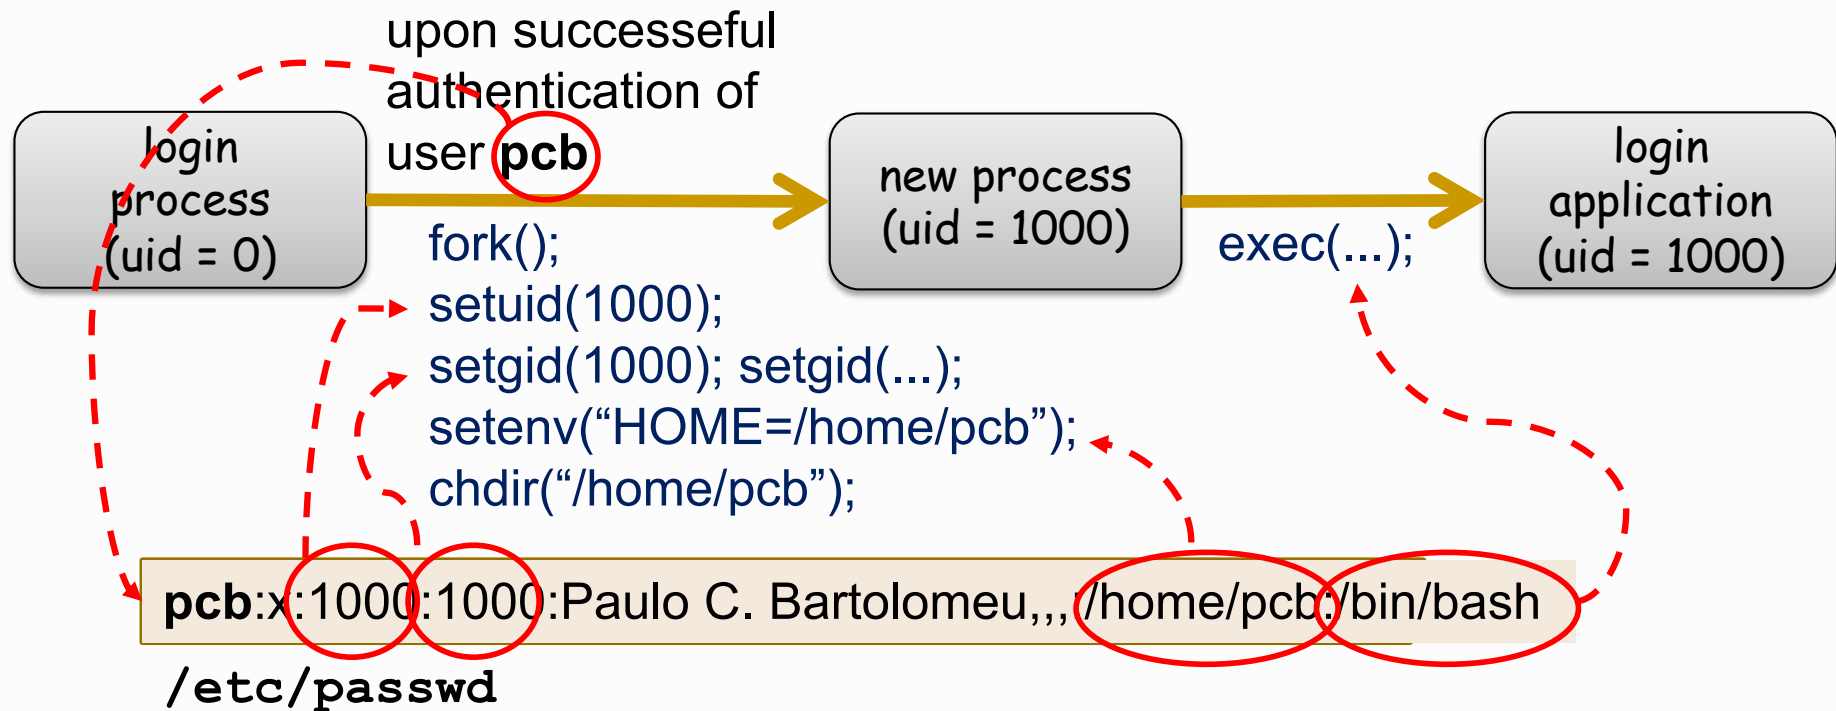
Linux login:

Not an OS kernel operation

- ▷ A privileged login application presents an interface for getting users' credentials
 - ♦ A username/password pair
 - ♦ Biometric data
 - ♦ Smartcard and activation PIN
- ▷ The login application validates the credentials and fetches the appropriate UID and GIDs for the user
 - ♦ And starts an initial user application on a process with those identifiers
 - In a Linux console this application is a shell (sh, bash, csh, tcsh, zsh, etc.)
 - ♦ When this process ends the login application reappears
- ▷ Thereafter all processes created by the user have its identifiers
 - ♦ Inherited through forks

Linux: from login to session processes

- ▶ The login process must be a privileged process
 - ♦ Has to create processes with arbitrary UID and GIDs
 - The ones of the entity logging in



Login in Linux:

Password validation process

- ▷ Username is used to fetch a UID/GID pair from `/etc/passwd`
 - ◆ And a set of additional GIDs in the `/etc/group` file
- ▷ Supplied password is transformed using a digest function
 - ◆ Currently configurable, for creating a new user (`/etc/login.defs`)
 - ◆ Its identification is stored along with the transformed password
- ▷ The result is checked against a value stored in `/etc/shadow`
 - ◆ Indexed again by the username
 - ◆ If they match, the user was correctly authenticated
- ▷ File protections
 - ◆ `/etc/passwd` and `/etc/group` can be read by anyone
 - This is fundamental, for instance, for listing directories (why?)
 - ◆ `/etc/shadow` can only be read by root
 - Protection against dictionary attacks

Virtualization on Intel processors

Modes of operation

- real-address mode
 - mode used on power up or after a hard or soft reset
 - 16-bit registers
 - segmented memory space ($16 \times \text{segment} + \text{offset}$)
- protected mode (rings 0 to 3, Intel calls then protection levels)
 - virtual memory, paging, multitasking, 32-bits, 64-bits
- hypervisor mode (ring -1)
- system management mode ("god" mode, ring -2)
- management engine (ring -3, see <http://www.isci-conf.org/share/doc/InvitedTalkFengweiZhang.pdf>)

Changing the mode of operation

- CLP (current protection level)
- system calls
- interrupts
- gates
 - call gates
 - trap gates
 - interrupt gates
 - task gates
- **vmxon**, **vmxoff**, **vmlaunch**, **vmenter**, **vmexit**, **vmresume** (virtual machine control structure)

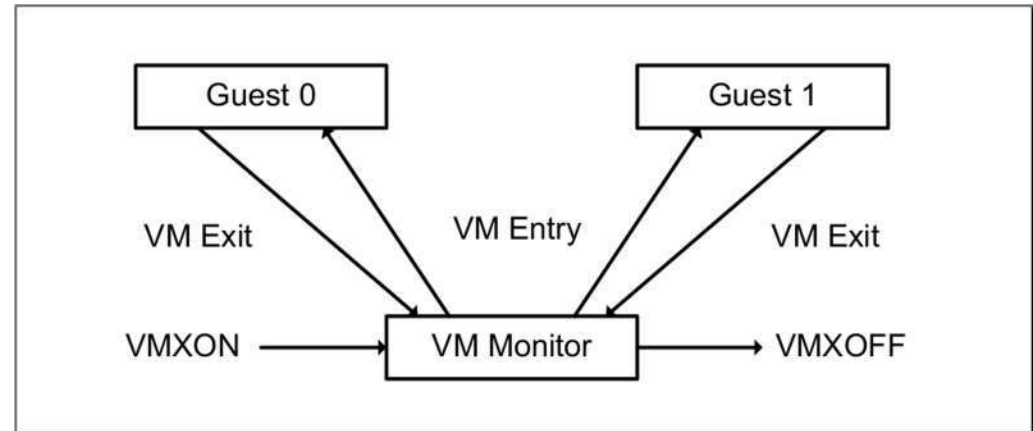


Figure 23-1. Interaction of a Virtual-Machine Monitor and Guests

Images source: Intel® 64 and IA-32 Architectures
Software Developer's Manual
Volume 3 (3A, 3B, 3C & 3D):
System Programming Guide

Virtual memory (paging, 3 levels, 4KiB pages)

- Translates virtual addresses to physical addresses

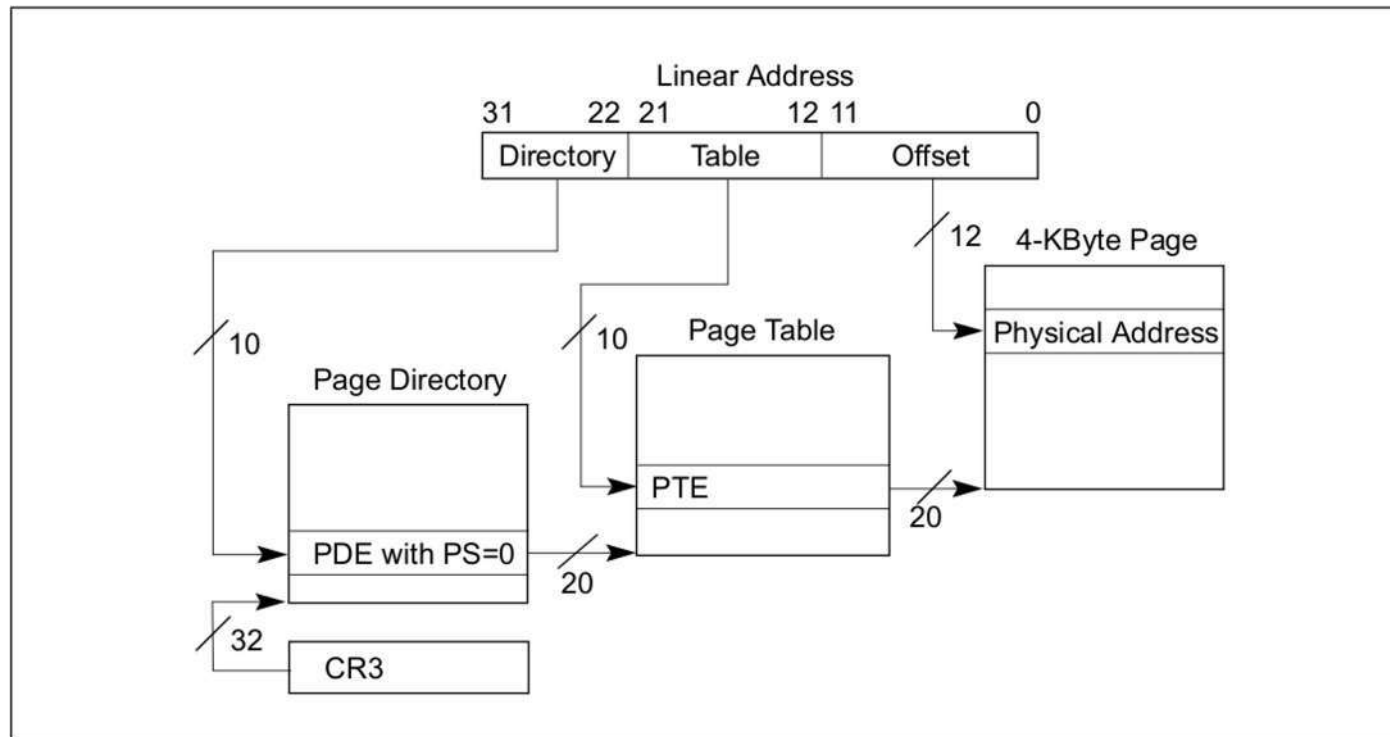


Figure 4-2. Linear-Address Translation to a 4-KByte Page using 32-Bit Paging

Virtual memory (paging, 2 levels, 4MiB pages)

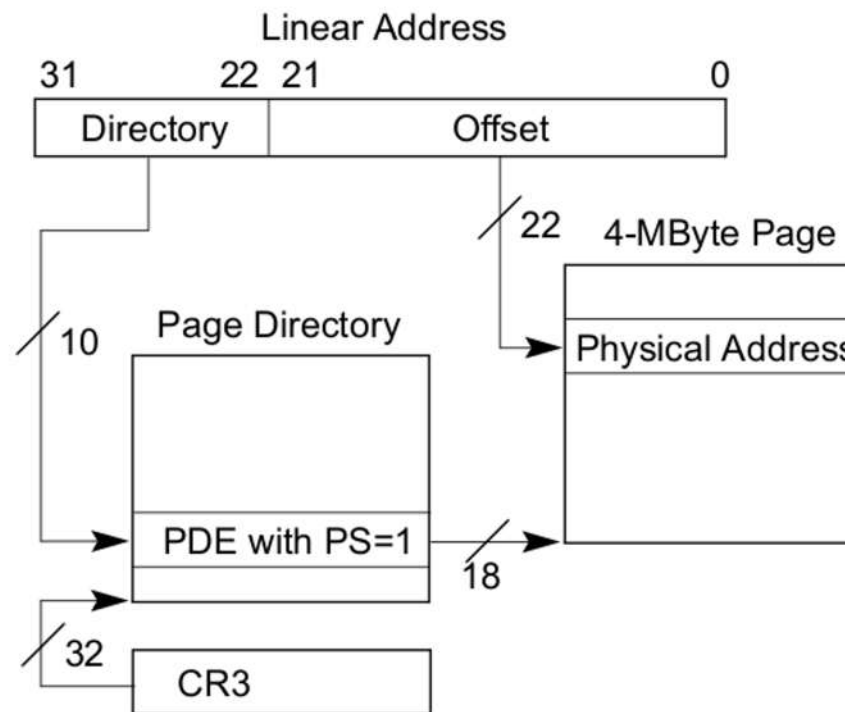


Figure 4-3. Linear-Address Translation to a 4-MByte Page using 32-Bit Paging

Virtual memory (4 levels, 4KiB pages)

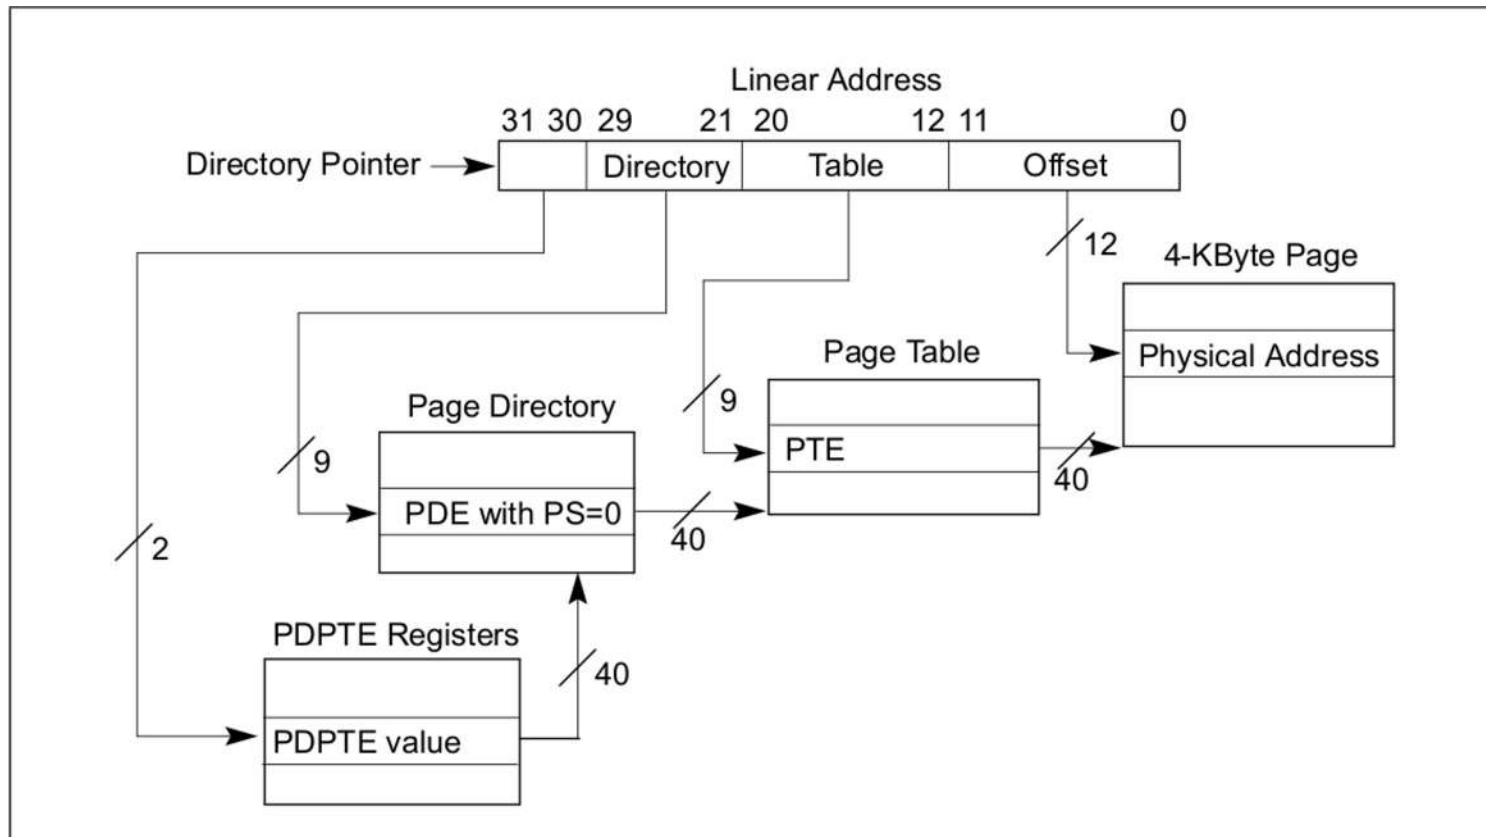


Figure 4-5. Linear-Address Translation to a 4-KByte Page using PAE Paging

Virtual memory (3 levels, 2MiB pages)

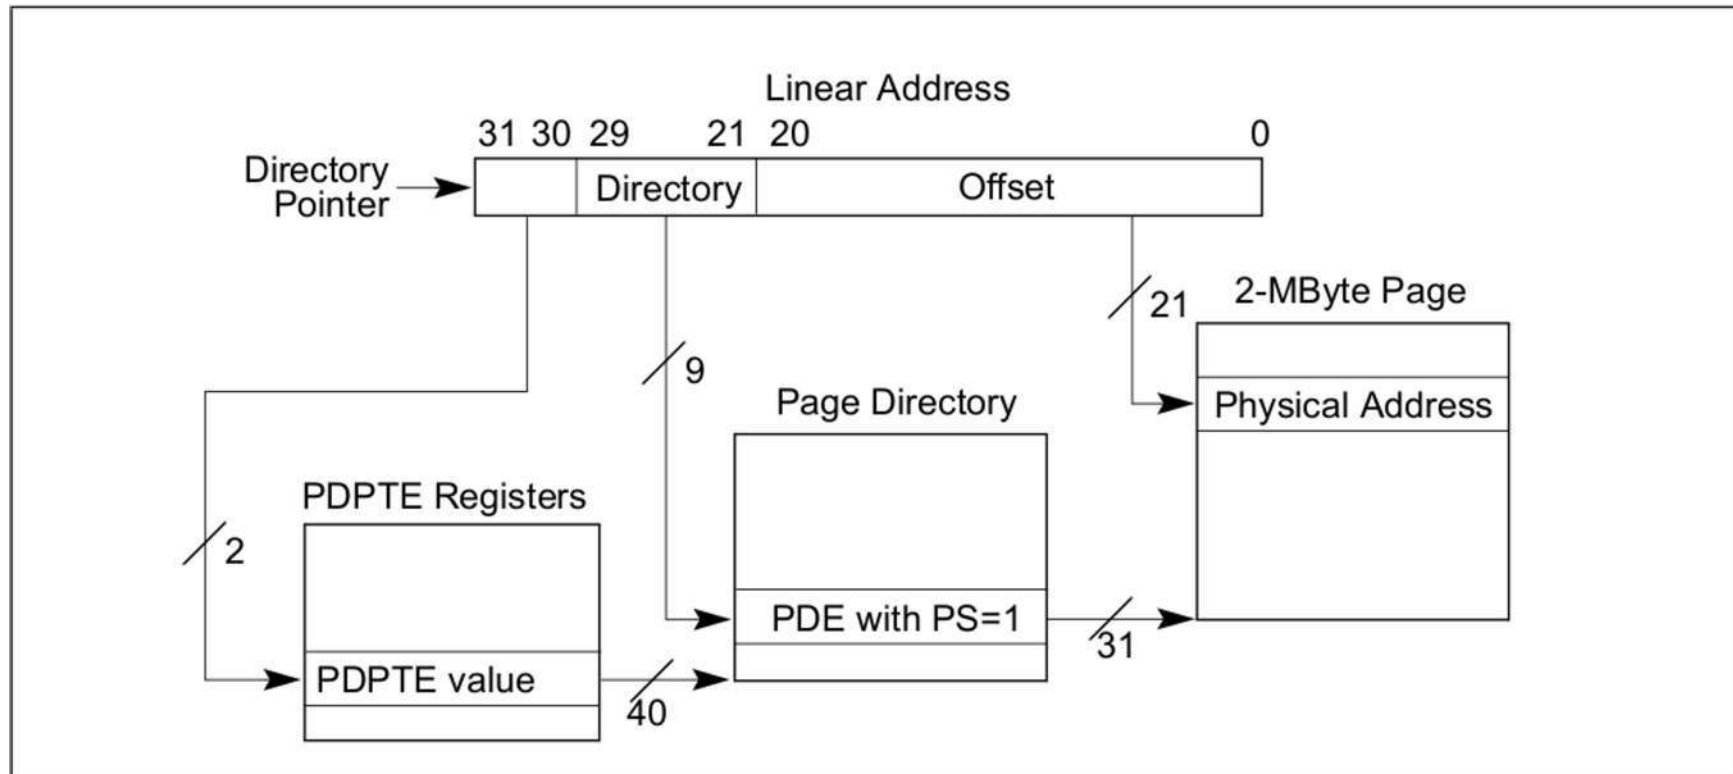


Figure 4-6. Linear-Address Translation to a 2-MByte Page using PAE Paging

Virtual memory (4 level, 4KiB pages)

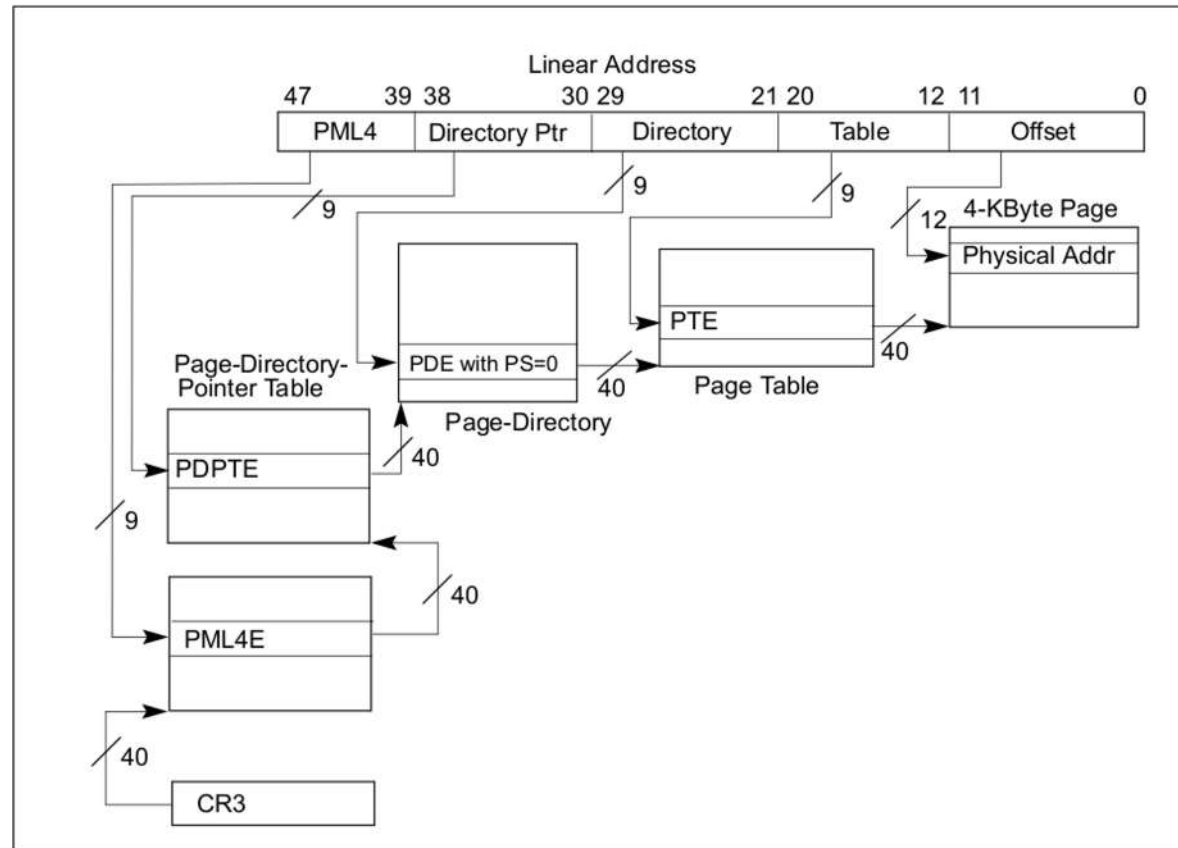


Figure 4-8. Linear-Address Translation to a 4-KByte Page using 4-Level Paging

How to put assembly instructions inside C code

Only C or only assembly

C code:

```
int f(int x)
{
    return 23 * x + 9;
}
```

Corresponding assembly code (Intel processor, AT&T syntax)

```
        .text
        .globl  f
        .p2align 4,,15
f:      imull    $23,%edi,%eax
        addl    $9,%eax
        ret
```

AT&T syntax (for a 64-bit Intel/AMD processor)

- Register names
 - 8-bit %al,%ah,%bl,%bh, ..., r8b, ..., r15b
 - 16-bit %ax,%bx, ..., r8w, ..., r15w
 - 32-bit %eax,%ebx,..., r8d, ..., r15d
 - 64-bit %rax, %rbx, ..., r8, ..., r15
- Constants are preceded by a dollar sign (\$)
- Destination at the end
- Instruction names include the register size at the end
 - 8 bits (char) b --- addb
 - 16 bits (short) w --- addw
 - 32 bits (int) l --- addl
 - 64 bits (long) q --- addq

Useful links

- <https://software.intel.com/content/www/us/en/develop/articles/intel-sdm.html>
- <https://developer.amd.com/resources/developer-guides-manuals/>
- <https://www.agner.org/optimize/#manuals>
- <https://gitlab.com/x86-psABIs/x86-64-ABI>
- <https://www.felixcloutier.com/x86/index.html>

How to put assembly instructions inside C code

C code with an embedded assembly instruction:

```
unsigned long first_one_bit(unsigned long word)
{
    register unsigned long result;

    asm("bsfq %[data], %[result]"
        : [result] "=r" (result)
        : [data] "r" (word)
        );
    return result;
}
```

How to put assembly instructions inside C code

A more complex example:

```
unsigned long read_time_stamp_counter(void)
{ // read tsc register (MUST be compiled with -m64)
  unsigned long tmp;

  asm volatile
  (
    "rdtsc ; "
    "shlq $32,%%rdx ; "
    "orq  %%rdx,%%rax"
    : "=a" (tmp)    /* output operands (a=rax register) */
    :               /* no input operands                */
    : "rdx","cc"    /* things that got modified                */
  );
  return tmp;
}
```

The gcc assembly instructions template

To insert assembly instructions inside a C function use the **asm** keyword as follows:

```
asm [volatile]
(
    assembler_template
    : output_operands
    : input_operands
    : clobbers
) ;
```

Useful links

- <https://gcc.gnu.org/onlinedocs/gcc/Extended-Asm.html>

The **volatile** keyword tells the compiler that the assembly code should not be moved (otherwise, during the optimization phase the compiler may place it in an unintended place).

Details (part 1)

- The **assembler_template** is a string containing the assembly source code
 - a pattern of the form %%reg refer to the specific register reg
 - a pattern of the form %[name] refers to a register holding one input or output argument (the compiler chooses the register that will be used)
- **output_operands** is a comma-separated list, possibly empty, of output or input/output parameters
 - each output parameter has the form [name] **constraint_string** (lvalue) where **constraint_string** can be (incomplete list):
 - "=r", meaning that the output is stored in a register (the register can be used in an unrelated input)
 - "&r", meaning that the output is stored in a register (the register can not be used as an input, early clobber)
 - "+r", meaning the argument is used as input and output, stored in a register

Details (part 2)

- **input_operands** is a comma-separated list, possibly empty, of input parameters
 - each input parameter has the form [name] **constraint_string** (C_expression), where **constraint_string** can be (among other possibilities)
 - "r", meaning that the input is stored in a register
 - "m", meaning that the input is stored in a memory position
 - note that input-only operands MUST NOT be modified by the assembly code
- **clobbers** is a comma-separated list, possibly empty, of things changed by the assembly code; these include specific register names, "cc" and "memory"

Details (part 3)

- It is possible to specify part of a register name using an extension of the %[name]; in particular
 - %b[name] specifies the low byte register name (bits 7..0)
 - %h[name] specifies the high byte register name (bits 15..8)
 - %w[name] specifies the low word register name (bits 15..0)
 - %k[name] specifies the low double word register name (bits 31..0)
 - unfortunately, the letter l is used for labels
 - %q[name] specifies the quad word register name (bits 63..0)
- Register usage conventions:
 - rbx, rbp and r12-r15 need to be saved if they are used
 - return value in rax
 - first 6 integer arguments in rdi, rsi, rdx, rcx, r8, and r9

A more elaborate example

```
void add3(unsigned long a[3], unsigned long b[3])
{ // a += b
  register unsigned long tmp;

  asm volatile(
    "movq\t(%[b]), %[tmp]"
    "\n\taddq\t%[tmp], (%[a])"
    "\n\tmovq\t8(%[a]), %[tmp]"
    "\n\tadcq\t%[tmp], 8(%[a])"
    "\n\tmovq\t16(%[a]), %[tmp]"
    "\n\tadcq\t%[tmp], 16(%[a])"
    : [tmp] "=&r" (tmp)
    : [a] "r" (a),
      [b] "r" (b)
    : "cc", "memory"
  );
}
```

```
.text
.p2align 4,,15
.globl add3
.type add3, @function

add3:
.LFB0:
.cfi_startproc
#APP
# 5 "z.c" 1
    movq    (%rsi), %rax
    addq    %rax, (%rdi)
    movq    8(%rdi), %rax
    adcq    %rax, 8(%rdi)
    movq    16(%rdi), %rax
    adcq    %rax, 16(%rdi)
# 0 "" 2
#NO_APP
    ret
.cfi_endproc
.LFE0:
.size add3, .-add3
.ident "GCC: (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0"
.section .note.GNU-stack,"",@progbits
```

Useful assembly instructions (part 1)

- the instruction **rdtsc** can be used to read the time stamp counter
 - it is a non-serializing instruction (the processor may reorder its execution)
 - the value returned depends on the core where it was executed
- the instruction **rdtscp** can be used to read the time stamp counter and the core signature (usually the code id number)
 - it is a serializing instruction
 - the value returned depends on the core where it was executed
 - `asm volatile("rdtscp" : "=a" (rax), "=c" (rcx), "=d" (rdx)) ;`
 - as in the **rdtsc** instruction, the counter value is given by $(rdx \ll 32) + rax$
 - the core signature is given by **rcx**; on GNU/Linux, its value is the core id that executed the instruction

Useful assembly instructions (part 2)

- the instruction `cplid` can be used to get information about the processor
 - it is a serializing instruction
 - this instruction unconditionally generates a trap (vmexit) when executed in a virtualized environment
 - it can be used when CPL is 3 (least privileged mode)
- the instruction `invd` can be used to invalidate a cache line
 - this instruction unconditionally generates a trap when executed in a virtualized environment
 - it can be used only when CPL is 0 (kernel mode)
 - use it only on a memory region whose contents are irrelevant

Useful links

- <https://en.wikipedia.org/wiki/CPUID>
- <https://software.intel.com/content/www/us/en/develop/articles/intel-sdm.html>

Useful assembly instructions (part 3)

- the instructions **rdseed** and **rdrand** are used to generate random numbers on recent Intel/AMD processors
 - the virtual machine hypervisor can set things up so that these instructions generate a trap
 - it can be used when CPL is 3 (least privileged mode)
- the instructions **in**, **out**, **rdtsc**, **rdmsrd**, and **rdpmcd** can also be used to generate a trap in a virtualized environment (if the hypervisors wants that to happen)



Intel Software Guard Extensions

What is SGX (Software Guard eXtensions)?

- It is a TEE (Trusted Execution Environment).
- Everything outside the processor chip is not trusted.
- In particular, the BIOS (Basic Input/Output System), the SMM (System Management Mode, ring -2), the ME (Intel Management Engine, ring -3), and the OS (Operating System, ring 0) are **not** trusted.
- The SGX code and data is put inside a special container (a SGX enclave).
- The contents of the enclave are signed (they are loaded from an untrusted source...) and can be attested by an external third party.
- The contents of the enclave are **isolated** from the rest of the system.
- The enclave code runs in **ring 3** (least privileged mode).
- All SGX instructions are implemented in microcode. The microcode is an integral part of the trust model (it is trusted).

SGX Enclave Memory

- As mentioned in the previous slide, the trust boundary perimeter is the processor chip (core, cache, and memory controller).
- So, the memory of the SGX enclave, when it resides outside of the processor chip (DRAM) is also encrypted.
- The memory encryption key is chosen at random after every processor reset.
- Values read from memory are checked to see if they match what was written (if not the processor hangs).
- This is done on a cache-line granularity (64 bytes) using a memory integrity tree.
 - For details, see <https://eprint.iacr.org/2016/204.pdf>
- Very small performance penalty if the SGX enclave memory footprint fits in the processor caches.

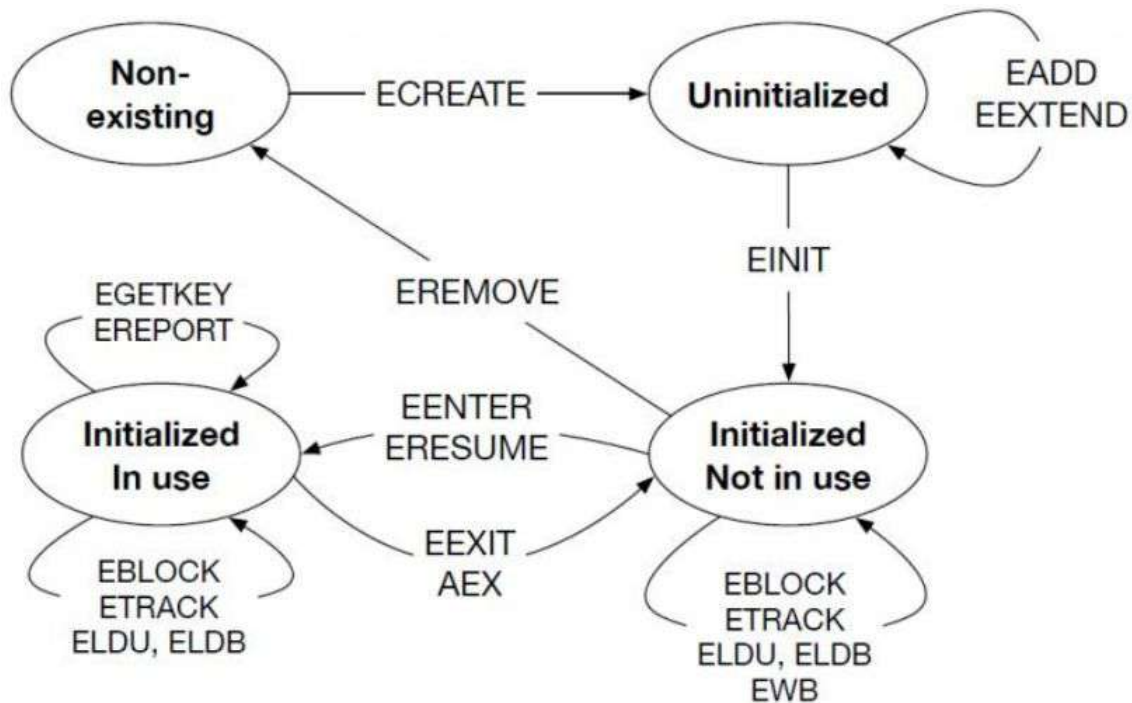
Instructions related to enclaves

- Ring 0 instructions
 - ECREATE, EADD and EINIT are used for Enclave Page Cache (EPC) management --- executed by privileged software such as an OS or a VMM
 - The EPC is an untrusted secure storage area used by the enclave; each 4KiB page has some security attributes that are stored in the Enclave Page Cache Map (EPCM), which is not accessible by software
- Ring 3 instructions
 - EENTER, EEXIT, EGETKEY, EREPORT and ERESUME are used by the user space software to execute functionality within or between enclaves.
- Illegal instructions inside an enclave
 - **cpuid**, **rdtsc**, input and output instructions, and some others are not allowed
 - **rdrand**/**rdseed** are allowed and can be virtualized (!?!)

Life cycle of an enclave

<https://software.intel.com/content/dam/develop/external/us/en/documents/intelsgx/enclavelifecycle.pdf>

1. Creation (ECREATE)
2. Loading (EADD, EEXTEND)
3. Initialization (EINIT)
4. Enter/Exit the Enclave (EENTER/EEXIT)
5. Teardown (EREMOVE)



Intel SGX Toolkit (version 2.13) requirements

<https://github.com/intel/linux-sgx>

- Hardware:
 - Intel 6th generation core processor or newer (but **not** some 10th or all 11th generation desktop processors!)
- 64-bit operating system:
 - Ubuntu 16.04, 18.04 or 20.04 LTS
 - Red Hat 7.6 or 8.2
 - CentOS 8.2
 - Fedora 31
- BIOS support (enabling SGX will reserve up to 128MiB of memory for the exclusive use of SGX enclaves)
- It's also possible to install it on Windows 10 (not covered in these slides)

Intel SGX Toolkit (version 2.13)

- Toolkit components:
 - Intel SGX kernel driver
 - Intel SGX PSW (Platform Software Package)
 - Intel SGX SDK
- Programming languages: **C** and **C++**
- Does my processor and OS support SGX (after BIOS configuration)?
 - `cpuid -1 | grep SGX`
 - If yes:
 - `SGX: Software Guard Extensions supported = true`
 - `SGX_LC: SGX launch config supported = true`

Intel SGX linux driver installation

- Install needed packages:
 - `sudo apt install build-essential ocaml automake autoconf libtool wget python3 libssl-dev dkms`
- Download driver (<https://01.org/intel-software-guard-extensions/downloads>)
 - `wget`
https://download.01.org/intel-sgx/sgx-linux/2.13/distro/ubuntu20.04-server/sgx_linux_x64_sdk_2.13.100.4.bin
- Install the Dynamic kernel Module Support (DKMS) driver:
 - `sudo bash sgx_linux_x64_driver_1.41.bin`
- If you are using secure boot, the kernel module has to be signed, and so this requires generating a new Machine-Owner Key (MOK). Just follow the instructions (a reboot will be required)
- the module location is `/lib/modules/5.8.0-48-generic/updates/dkms/intel_sgx.ko` and the module name is (obviously) `intel_sgx`.

Intel SGX PSW installation (on ubuntu)

- Install needed packages:
 - `sudo apt install libssl-dev libcurl4-openssl-dev libprotobuf-dev`
- Run the following commands
 - `echo 'deb [arch=amd64]
https://download.01.org/intel-sgx/sgx_repo/ubuntu focal main' |
sudo tee /etc/apt/sources.list.d/intel-sgx.list`
 - `wget -qO -
https://download.01.org/intel-sgx/sgx_repo/ubuntu/intel-sgx-deb.ke
y | sudo apt-key add -`
 - `sudo apt update`
 - `sudo apt install libsgx-launch libsgx-urts`
 - `sudo apt install libsgx-epid libsgx-urts`

Intel SGX SDK installation (on ubuntu)

- Do the following:
 - `wget https://download.01.org/intel-sgx/latest/linux-latest/distro/ubuntu20.04-server/sgx_linux_x64_sdk_2.13.100.4.bin`
 - `sudo bash sgx_linux_x64_sdk_2.13.100.4.bin`
 - answer no and choose `/opt/intel` as the installation directory
 - copy the contents of `/opt/intel/sgxsdk/environment` to your `.bashrc`
 - `wget https://download.01.org/intel-sgx/latest/linux-latest/as.ld.objdump.gold.r3.tar.gz`
 - `tar xzvf as.ld.objdump.gold.r3.tar.gz external/toolset/ubuntu20.04`
 - `sudo cp -v external/toolset/ubuntu20.04/* /usr/local/bin/`

Intel SGX SDK test (on ubuntu)

- Do the following:
 - `mkdir tmp`
 - `cd tmp`
 - `cp -av /opt/intel/sgxsdk/SampleCode/SampleEnclave .`
 - `cd SampleEnclave`
 - `make SGX_DEBUG=0 SGX_PRERELEASE=1`
 - `./app`
 - `make clean`
- The output should be

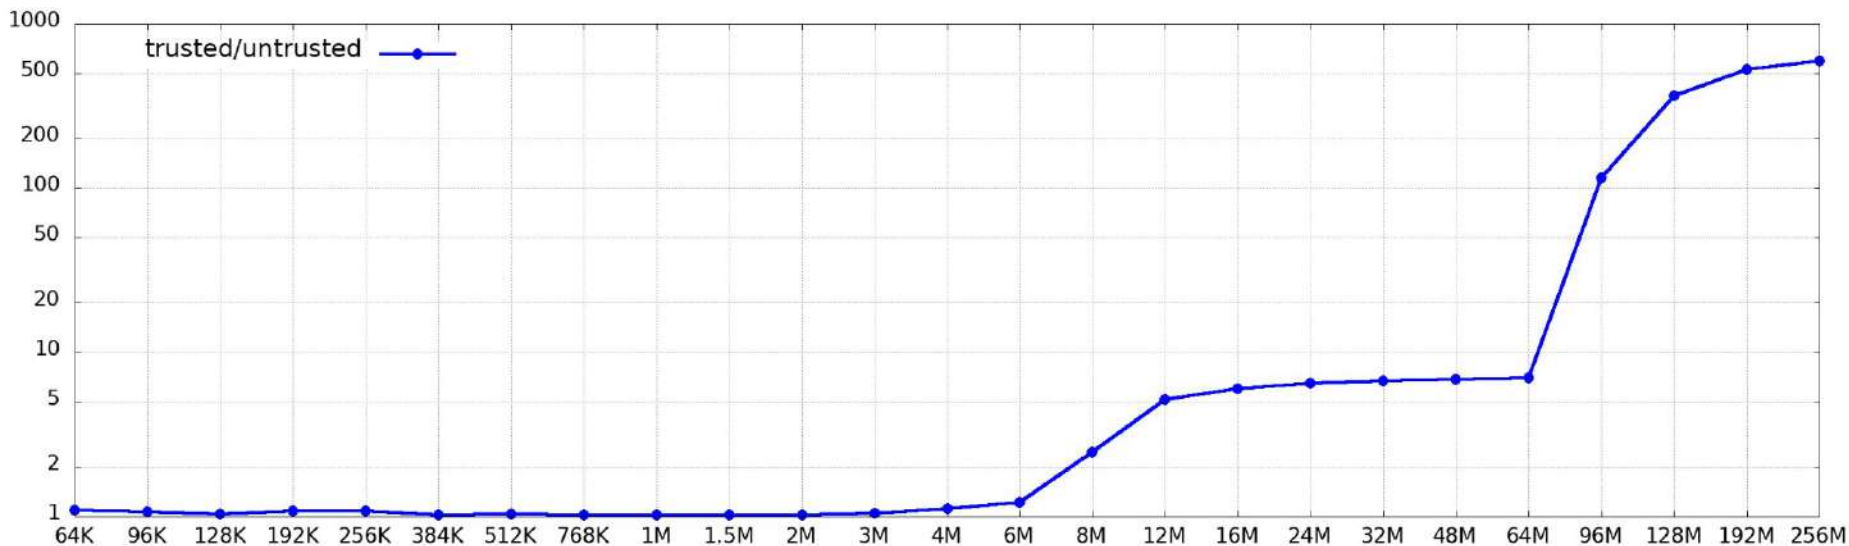
```
Checksum(0x0x7ffeac1ee4f0, 100) = 0xfffd4143
Info: executing thread synchronization, please wait...
Info: SampleEnclave successfully returned.
Enter a character before exit ...
```

Guidelines for designing applications using SGX

- Partition the software into trusted and untrusted components
- Use the SGX SDK tools to create the enclave module (a shared object) --- it implements the trusted component of the software
- The enclave code and data **is not secret**
- Secrets has to be loaded in a secure manner (using an ECDH key exchange for example) from a trusted outside source.
- Enclave data has to be **sealed** (encrypted and signed) if it is stored outside of the enclave
- Enclave data has to be **unsealed** if it is loaded into the enclave
- To minimize side-channels attack vectors, keep the secrets in the sealed state, and only unseal them for as long as it is truly necessary

Performance Overhead

- Creating an enclave is relatively slow (4KiB pages have to be added one at a time)
- Calling an enclave function from outside of the enclave (an **ecall**) or calling a non enclave function from inside the enclave (an **ocall**) takes about 10k clock cycles



Sealing and unsealing

<https://software.intel.com/content/www/us/en/develop/blogs/introduction-to-intel-sgx-sealing.html>

- Each Intel processor has one unique key, set during manufacturing, not known to Intel (presumably).
- Each SGX enclave has an identity, which is a hash of the enclave log as it goes through every step of the build and initialization process.
- The enclave identity (value of **MRENCLAVE**), together with the processor's key, should therefore be truly unique.
- It should be used to seal and unseal data.
- Because the enclave identity changes when the enclave code is modified, migration of the sealed data to a new software version has to be done with extreme care.
- One possibility is to create an encrypted secure channel between the two software versions, and to migrate the data in that way.

SDK documentation

- [Developer reference for Linux OS](#) (PDF)
- [Intel developer zone --- Software Guard Extensions](#) (online)

Intel SGX SDK compilation modes

- SGX applications can be compiled in several modes:
 - hardware debug mode (signed with Intel's key, code not optimized)
 - `SGX_MODE=HW SGX_DEBUG=1 SGC_PRERELEASE=0`
 - hardware prerelease mode (signed with your key, code is optimized)
 - `SGX_MODE=HW SGX_DEBUG=0 SGC_PRERELEASE=1`
 - hardware release mode (signed with your key, code is optimized, cannot be debugged)
 - `SGX_MODE=HW SGX_DEBUG=0 SGC_PRERELEASE=0`
 - This mode may require a [commercial licence](#)
 - simulation mode (in debug mode)
 - `SGX_MODE=SIM SGX_DEBUG=1`

SGX SDK Tools

- Edger8r (**`sgx_edger8r`**)
 - Generates "edge" routines (interface between the untrusted application and the enclave) described in a Enclave Description Language (EDL) file
 - Using it on file XYZ.edl produces files XYZ_[tu]t.[hc] where t=trusted, u=untrusted, h=prototypes, and c=functions
- Enclave signing tool (**`sgx_sign`**)
 - supports key management
- Enclave Memory Measurement Tool (**`sgx_emmt`**)
 - Use it to measure how much memory the enclave uses (needed by the Enclave Configuration File)

Writing Enclave Functions

- Describe each function that may be called from outside of the enclave in the .edl file
- The functions can use special versions of the **C/C++** runtime libraries (available in the SDK)
- System calls are not allowed (use **ocalls** instead; **C** linkage only!)
- Not all **C/C++** language features are available
- The `sgx_edger8r` tool will take care of the details of making the execution flow enter or leave an enclave
 - Pointer arguments in **ecall** functions must point to untrusted memory
 - Pointer arguments in **ocall** functions must point to trusted memory
 - You may need to copy buffer from untrusted memory to trusted memory
- Keep in mind that the enclave will be statically linked

Some available trusted libraries

- `libsgx_tstdc.a` (standard **C** library, math, strings, etc.)
- `libsgx_tcxx.a` (standard **C++** libraries, STL)
- `libsgx_tservice.a` (seal/unseal, EC DH library, etc.)
- `libsgx_tcrypto.a`
- `libsgx_tkey_exchange.a`
- `libsgx_tpcl.a` (Protected Code Loader, for enclave code confidentiality)

Hello world in an enclave

(https://github.com/sangfansh/SGX101_sample_code)

- One enclave function, `printf_helloworld`, prints the text "Hello World"
- It cannot do this directly, so it calls an untrusted function, `ocall_printf_string`, to do the actual printing
- In this example, for the enclave code, the `printf` function is re-implemented so that its output goes to a string
- In the untrusted part, the enclave is loaded and the `printf_helloworld` function is called

Hello world in an enclave

- List of files
 - **Makefile**
 - **App/App.cpp**
 - **App/App.h**
 - **Enclave/Enclave.config.xml**
 - **Enclave/Enclave.cpp**
 - **Enclave/Enclave.edl**
 - **Enclave/Enclave.h**
 - **Enclave/Enclave.lds**
 - **Enclave/Enclave_private.pem**

- Relevant parts of the **Makefile**:
 - **SGX_SDK** **?= /opt/intel/sgxsdk**
 - **SGX_MODE** **?= HW**
 - **SGX_ARCH** **?= x64**
 - **SGX_DEBUG** **?= 1**
- You can also add
 - **SGX_PRERELEASE** **?= 0**
- List your untrusted source code files (the application) in the
App Settings
section
- List your trusted source code files (the SGX enclave) in the
Enclave Settings
section

- Relevant parts of `App/App.h`:

```
#include "sgx_error.h"    /* sgx_status_t */
#include "sgx_eid.h"      /* sgx_enclave_id_t */

#define TOKEN_FILENAME    "enclave.token"
#define ENCLAVE_FILENAME "enclave.signed.so"

extern sgx_enclave_id_t global_eid; /* global enclave id */
```

- Relevant parts of `App/App.cpp`:

```
int initialize_enclave(void) { /*...*/ }
```

```
void ocall_print_string(const char *str){ printf("%s",str); }
```

```
int SGX_CDECL main(int argc, char *argv[])  
{  
    if(initialize_enclave() < 0) return -1;  
    printf_helloworld(global_eid);  
    sgx_destroy_enclave(global_eid);  
}
```

- Enclave/Enclave_config.xml:

<EnclaveConfiguration>	
<ProdID>0</ProdID>	user defined
<ISVSVN>0</ISVSVN>	user defined
<StackMaxSize>0x40000</StackMaxSize>	stack size
<HeapMaxSize>0x100000</HeapMaxSize>	heap size
<TCSNum>10</TCSNum>	maximum number of threads
<TCSPolicy>1</TCSPolicy>	0=bound threads
<DisableDebug>0</DisableDebug>	0=enclave can be debugged
<MiscSelect>0</MiscSelect>	???
<MiscMask>0xFFFFFFFF</MiscMask>	???
</EnclaveConfiguration>	

- Relevant parts of `Enclave/Enclave.edl`:

```
enclave {  
    trusted {  
        public void printf_helloworld();  
    };  
    untrusted {  
        void ocall_print_string([in,string] const char *str);  
    };  
};
```

- Relevant parts of `Enclave/Enclave.h`:

```
#include <stdlib.h>
#include <assert.h>

#ifdef __cplusplus
extern "C" {
#endif
void printf(const char *fmt, ...);
void printf_helloworld();
#ifdef __cplusplus
}
#endif
```

- Relevant parts of `Enclave/Enclave.cpp`:

```
#include <stdarg.h>
#include <stdio.h>      /* vsnprintf    */
#include "Enclave.h"
#include "Enclave_t.h"  /* print_string */

void printf(const char *fmt, ...) { char buf[BUFSIZ];
    va_list ap; va_start(ap, fmt);
    vsnprintf(buf, BUFSIZ, fmt, ap);
    va_end(ap); ocall_print_string(buf); }

void printf_helloworld() { printf("Hello World\n"); }
```

Power events

- The enclave page cache is **destroyed** then the processor is put in state S2 or lower
 - S0 --- The CPU is executing instructions, and background tasks are running even if the system appears idle and the display is powered off.
 - S1 --- Processor caches are flushed, CPU stops executing instructions. Power to CPU and RAM is maintained. Devices may or may not power off. This is a high-power standby state, sometimes called “power on suspend.”
 - S2 --- CPU is powered off. CPU context and contents of the system cache are lost.
 - S3 --- RAM is powered on to preserve its contents. A standby or sleep state.
 - S4 --- RAM is saved to nonvolatile storage in a hibernation file before powering off. When powered on, the hibernation file is read in to restore the system state. A hibernation state.
 - S5 --- “Soft off.” The system is off but some components are powered to allow a full system power-on via some external event, such as Wake-on-LAN, a system management component, or a connected device.
- Follow the following [advice](#) in order to deal with power events.

Some extra useful links

- [Intel SGX explained](#)
- <https://sgx101.gitbook.io/sgx101/>
- [SGX Intel tutorial series](#)



deti

universidade de aveiro
departamento de eletrónica,
telecomunicações e informática

ARM TrustZone

SoC and IP

▷ SoC (System-on-Chip)

- ◆ Tackles the provisioning of complex and application-specific, multifunctional processors
- ◆ The major functional components of a complete end-product are integrated into a single chip

▷ Intellectual property (IP) modules

- ◆ Pre-designed, reusable electronic components for hardware chips

SoC structure

- ▶ A SoC usually contains
 - ♦ Processors
 - ♦ IPs
 - Namely security IPs
 - ♦ Memory elements (RAM, ROM, etc.)
 - ♦ Buses

ARM TrustZone

- ▶ Set of technologies for packing special security features into a SoC
 - ♦ Extra security-related features on processor cores
 - Instructions
 - Bus lines
 - Execution levels
 - Extra logic for dealing with interruptions
 - ♦ Security-related IPs

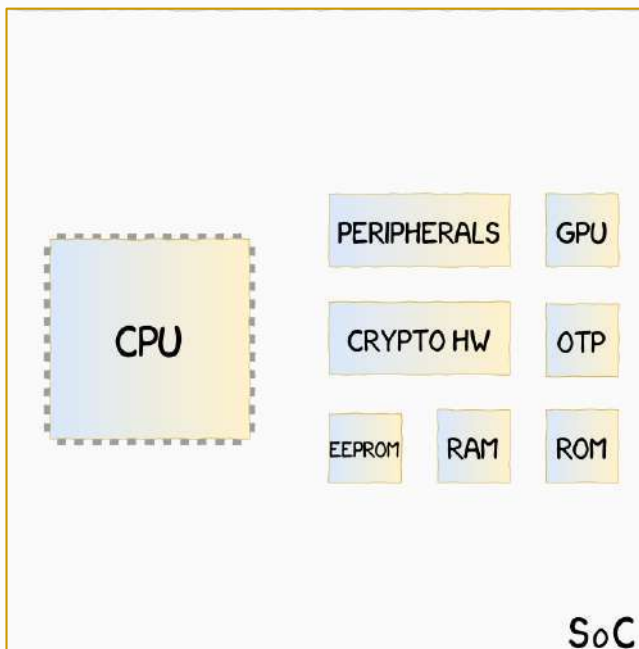
ARM TrustZone: goal

- ▷ TEE for ARM-powered embedded systems
 - ◆ Providing hardware-based isolation
- ▷ It allows to run a trusted system in parallel with the main operation system
 - ◆ Rich OS
 - Where most applications will run
 - ◆ Secure (or Trusted) OS
 - Where secure (or trusted) applications will run
 - It can be a simple library, and not a full-fledged OS

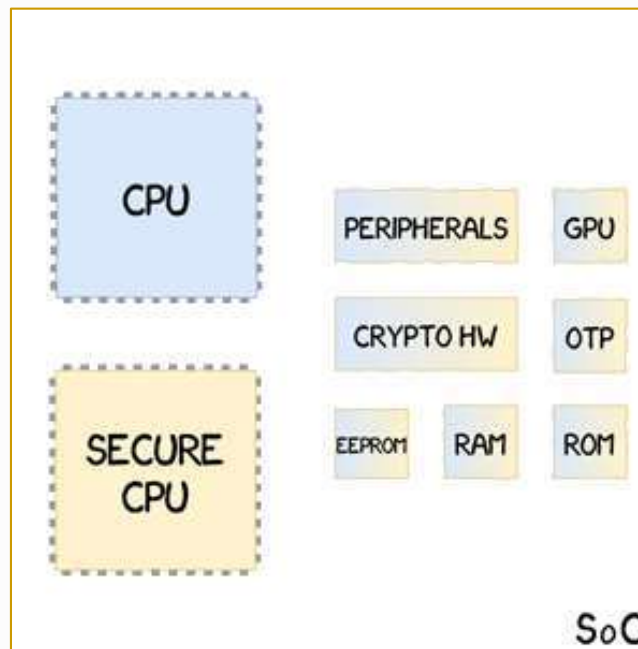
ARM TrustZone:

Comparison with other similar TEEs


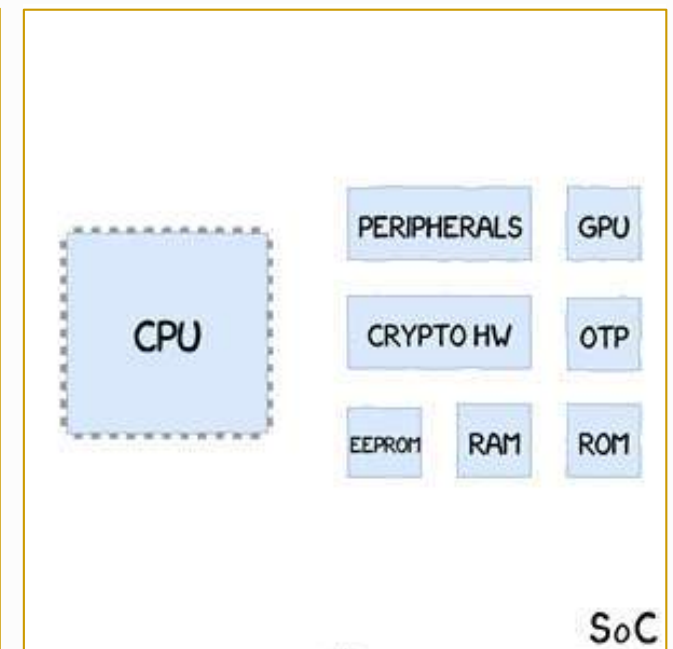
ARM TrustZone



Apple SEP



Google Titan M



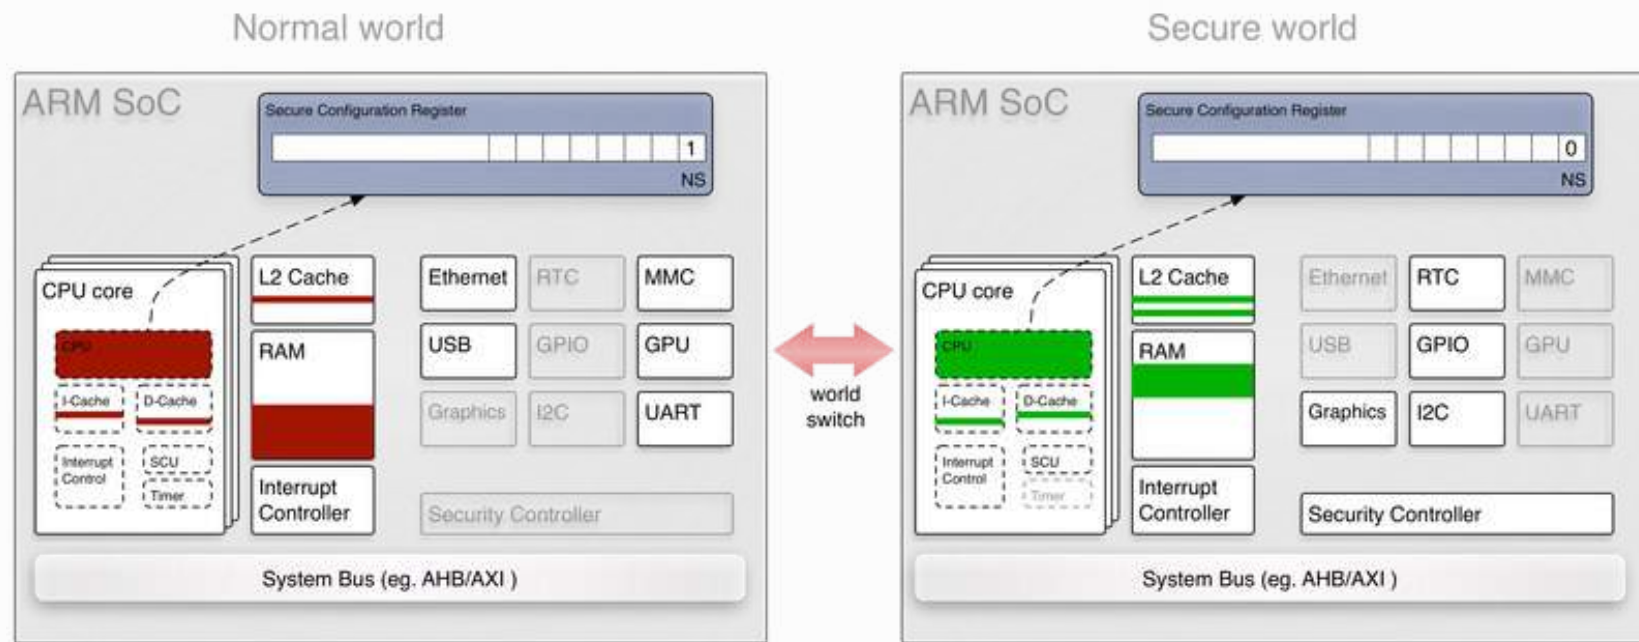
<https://blog.quarkslab.com/a-deep-dive-into-samsungs-trustzone-part-1.html>

Worlds

- ▶ Isolation is achieved by exploring the same CPU in two different worlds (or states)
 - ◆ Normal world → for running the Rich OS
 - ◆ Secure world → for running the Secure OS
- ▶ A CPU flag bit defines the current world
 - ◆ NS bit of the SCR (Secure Configuration Register)
 - ◆ 0 – Secure state
 - ◆ 1 – Non-secure state

Protected hardware resources

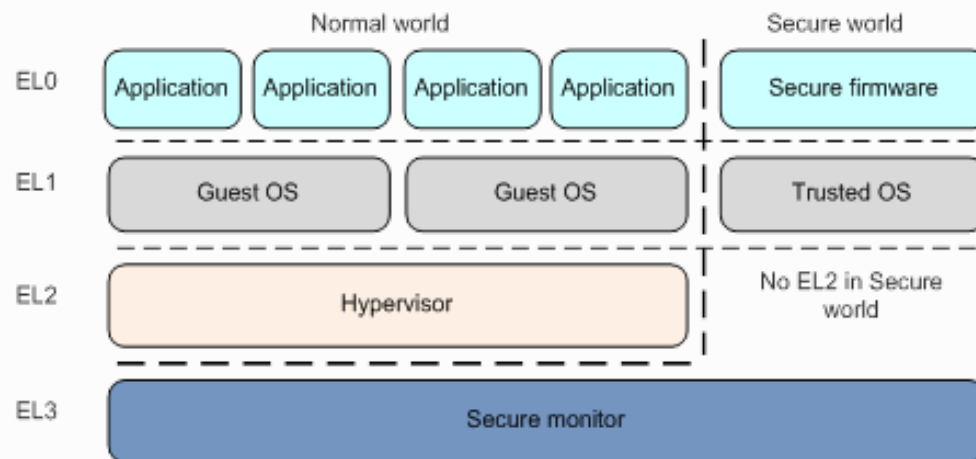
- Resources accessible only to the Secure OS



<https://genode.org/documentation/articles/trustzone>

ARM (v8) exception levels

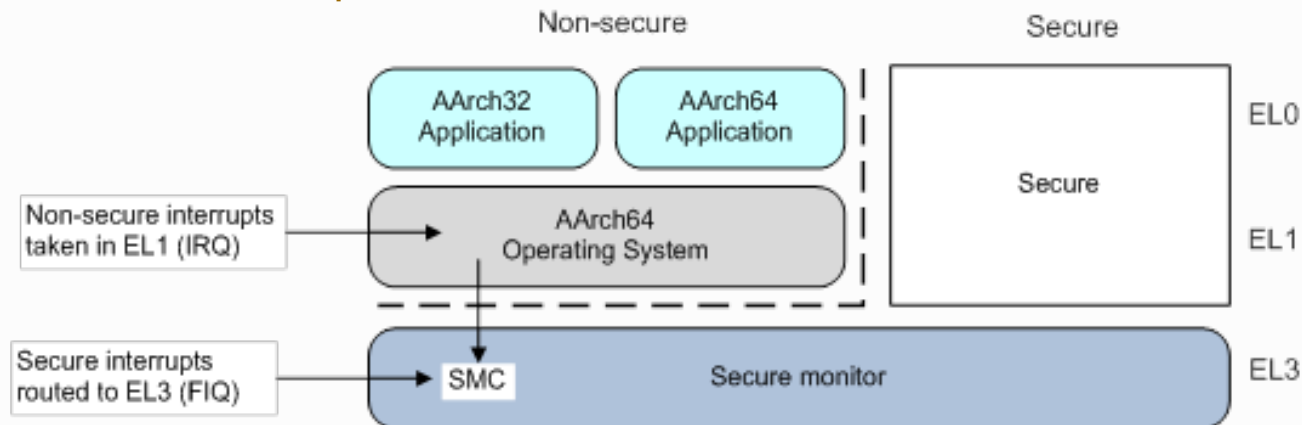
- ▷ Similar to run levels
- ▷ TrustZone introduces one EL more
 - ◆ Secure monitor (EL3)
- ▷ Combination of exception levels



<https://sergioprado.blog/introduction-to-trusted-execution-environment-tee-arm-trustzone/>

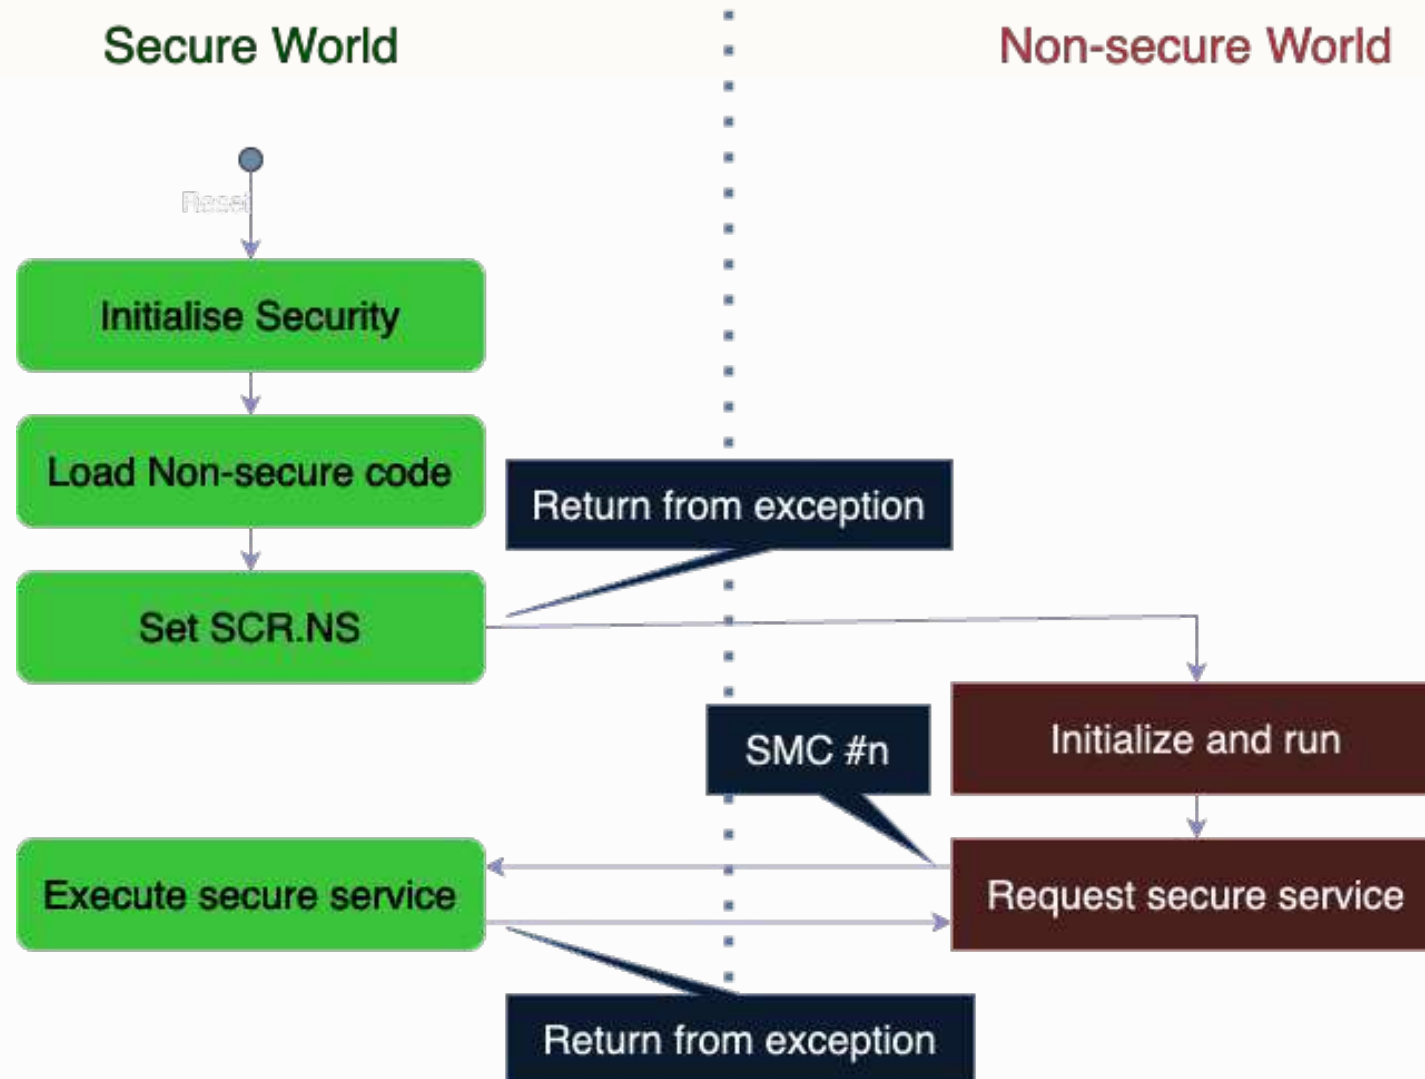
Access to the Secure world

- ▷ Calls from the Rich OS
 - ♦ SMC (Secure Monitor Call)
 - ♦ Typically implemented by Rich OS drivers
- ▷ Interrupts from the Secure hardware
 - ♦ Must be handled by the Secure OS
- ▷ Both enter first in EL3
 - ♦ Then are dispatched to the Secure world



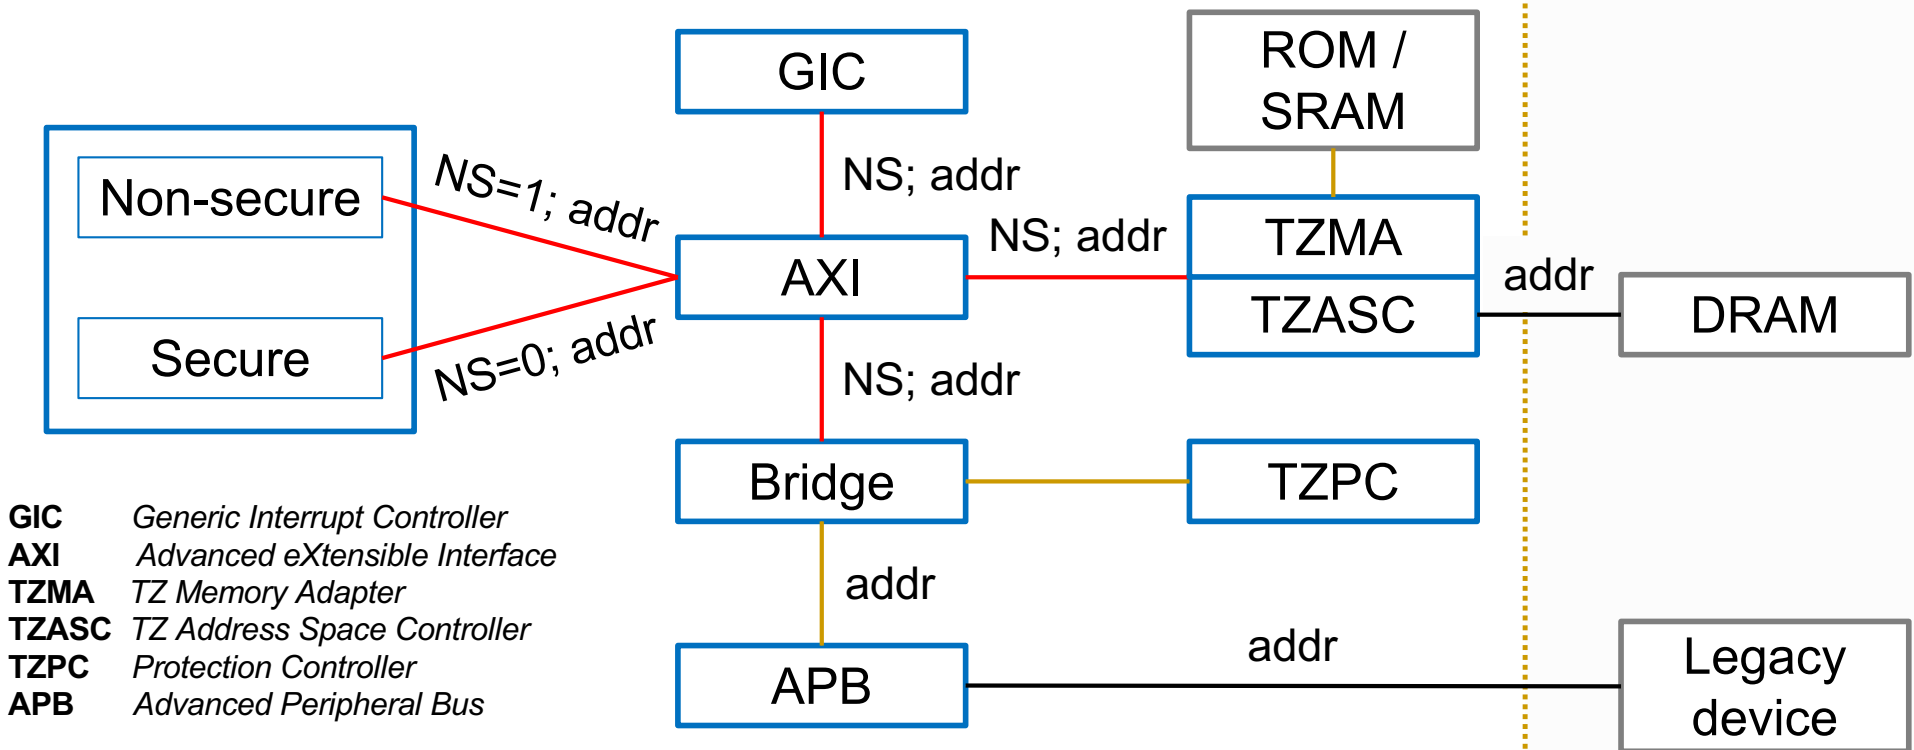
<https://sergioprado.blog/introduction-to-trusted-execution-environment-tee-arm-trustzone/>

Access to the Secure world



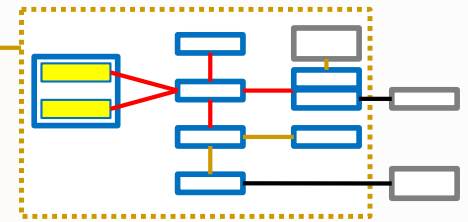
Architecture overview

SoC



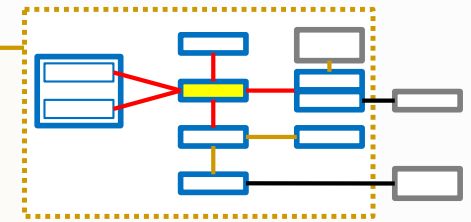
Architectural details:

MMU / TLB / Cache Controllers

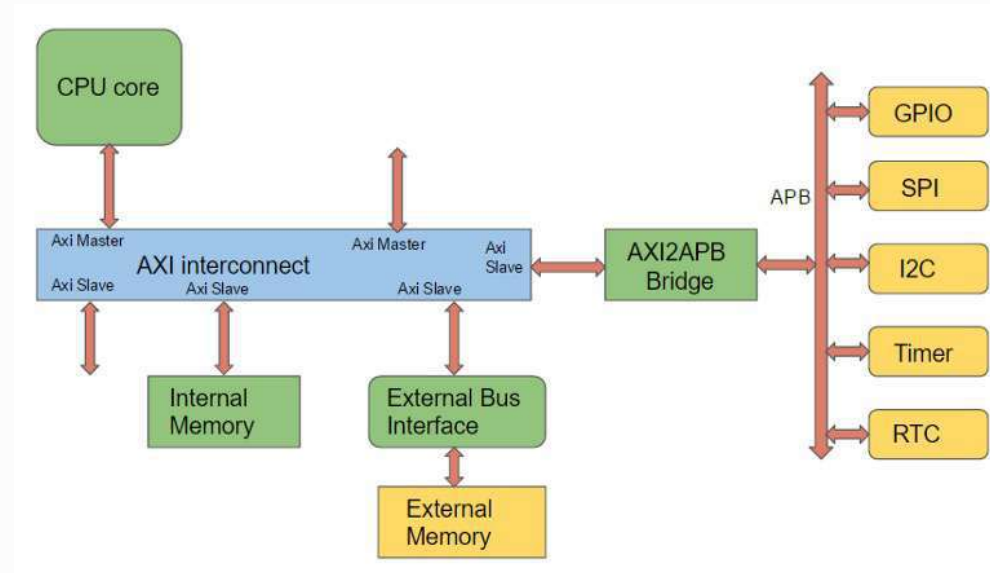


- ▷ 2 separate, virtual MMUs
 - ◆ Indexed by NS
- ▷ Single TLB
 - ◆ But entries keep the value of NS that created them
 - ◆ No need to invalidate them when switching between worlds
- ▷ The Secure world can still access non-secure memory
 - ◆ Extra bit on each entry in the secure translation table
- ▷ Single cache
 - ◆ Cache lines keep the NS address bit

Architectural details: AXI (Advanced eXtensible Interface)



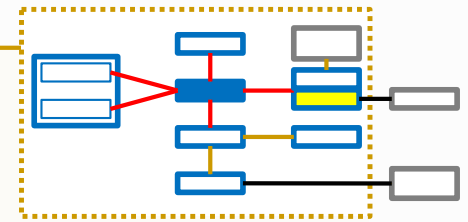
► SoC internal bus



- ## ► Extra NS line for secure read/write operations
- ◆ Non-secure master cannot access a resource marked as secure

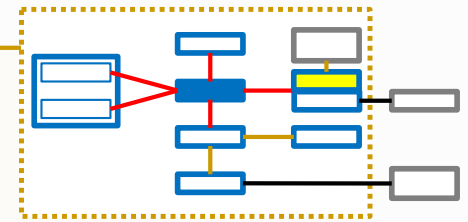
<https://anysilicon.com/understanding-amba-bus-architecture-protocols>

Architectural details: TZASC (TZ Address Space Controller)



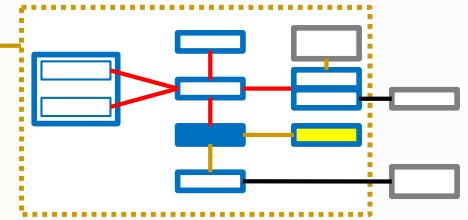
- ▶ Allows a dynamic classification of AXI slave memory-mapped devices as secure or non-secure
 - ◆ Partitioning of single memory units
- ▶ Controlled by the Secure world

Architectural details: TZMA (TZ Memory Adapter)



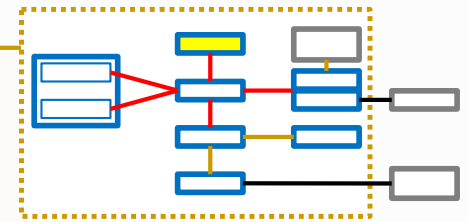
- ▶ Keeps a classification of in-SoC memory areas as secure and non-secure
 - ◆ ROM or SRAM
- ▶ Non-secure accesses cannot access secured memory areas
- ▶ Controlled by the Secure world

Architectural details: TZPC (TZ Protection Controller)



- ▶ Allows to dynamically set the security of a peripheral connected to the APB (Advanced Peripheral Bus)
 - ◆ Protects non-secure access requests to reach peripherals marked as secure
- ▶ Controlled by the Secure world

Architectural details: GIC (Generic Interrupt Controller)



- ▷ Classifies interrupts as secure or non-secure
 - ◆ Once set, cannot be changed
- ▷ Interrupts can be normal or fast (high-priority)
 - ◆ Secure interrupts usually have higher priority
- ▷ Interrupts with a security classification different from the current world force the switching to Monitor (EL3)
- ▷ Controlled by the Secure world

TrustZone bootstrap

- ▷ A TZ-enable ARM SoC boots on the secure world
 - ◆ It allows the Secure world to configure the TZ-related components to enforce a given security policy
- ▷ The configuration data can be
 - ◆ Embedded in the SoC ROM
 - ◆ Provided by external peripherals and validated with information in SoC ROM
 - e.g. must contain a signature validated with a in-SoC public key

Linux kernel namespaces

namespaces

- ♦ A **namespace wraps a global system resource in an abstraction** that makes it appear to the processes within the namespace that they have their own isolated instance of the global resource;
- ♦ Changes to the global resource are visible to other processes that are members of the namespace, but are invisible to other processes;
- ♦ One use of namespaces is to implement containers.

Linux Programmer's Manual: *\$man namespaces*

In a nutshell: restrictions and translations

- ▷ Extend the CHROOT concept to other resources
 - ◆ CHROOT enables each process to have a different notion of the file system root
- ▷ Restrictions
 - ◆ Limit the number of resources a process can use
- ▷ Translations
 - ◆ Access resources with a name different from the one the process thinks it is using

Advantages

▶ Enforce the Principle of Least Privilege

- ◆ Processes cannot make use of resources they don't need for the application they run
- ◆ If their application gets compromised, the processes will prevent their misbehavior to access forbidden resources

Process namespace

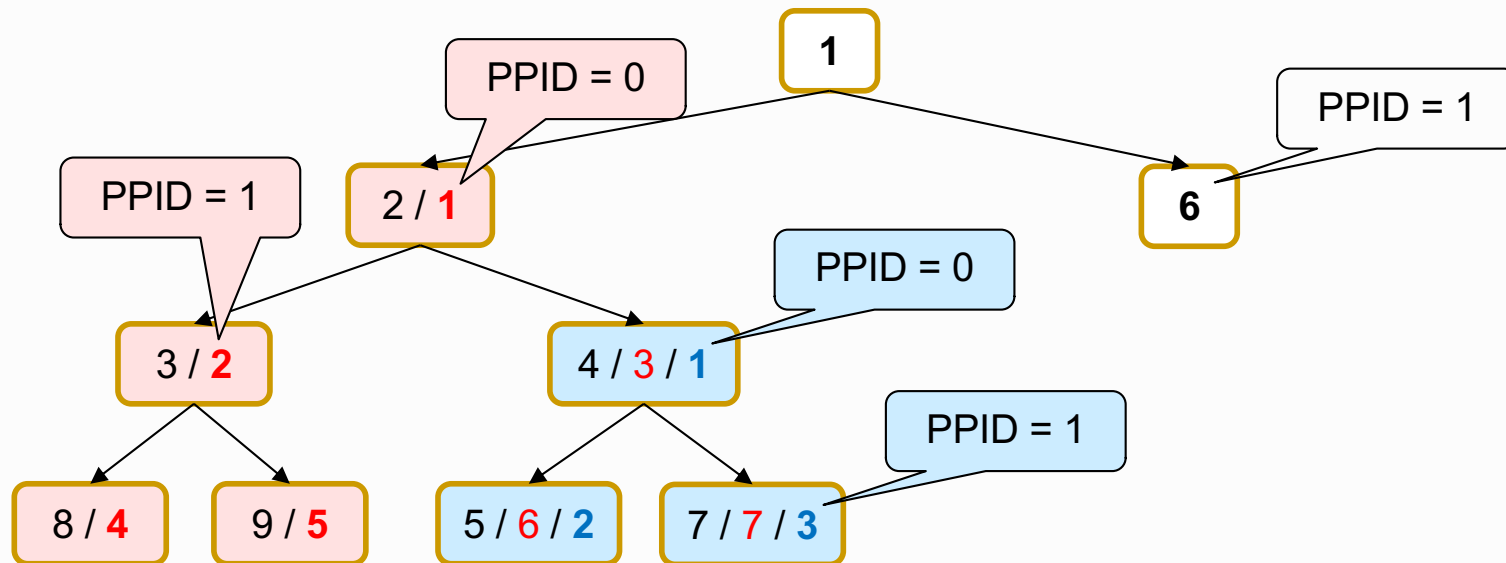
- ♦ Process namespaces **isolate the process ID number space**, meaning that processes in different PID namespaces can have the same PID;
- ♦ Process namespaces allow containers to provide functionality such as **suspending/resuming the set of processes** in the container and migrating the container to a new host while the processes inside the container maintain the same PIDs;
- ♦ PIDs in a **new PID namespace start at 1**, somewhat like a standalone system, and calls to fork or clone will produce processes with PIDs that are unique within the namespace.

Linux Programmer's Manual: *\$man pid_namespaces*

Process namespace

- ▶ UNIX processes form a parent-child tree
 - ◆ Process identifiers (PIDs) are global
- ▶ The process namespace creates nested trees
 - ◆ Processes within inner namespaces cannot access processes of outer namespaces by their PID
- ▶ A process can have many PIDs
 - ◆ One for each namespace it belongs

Process namespaces



Network namespace

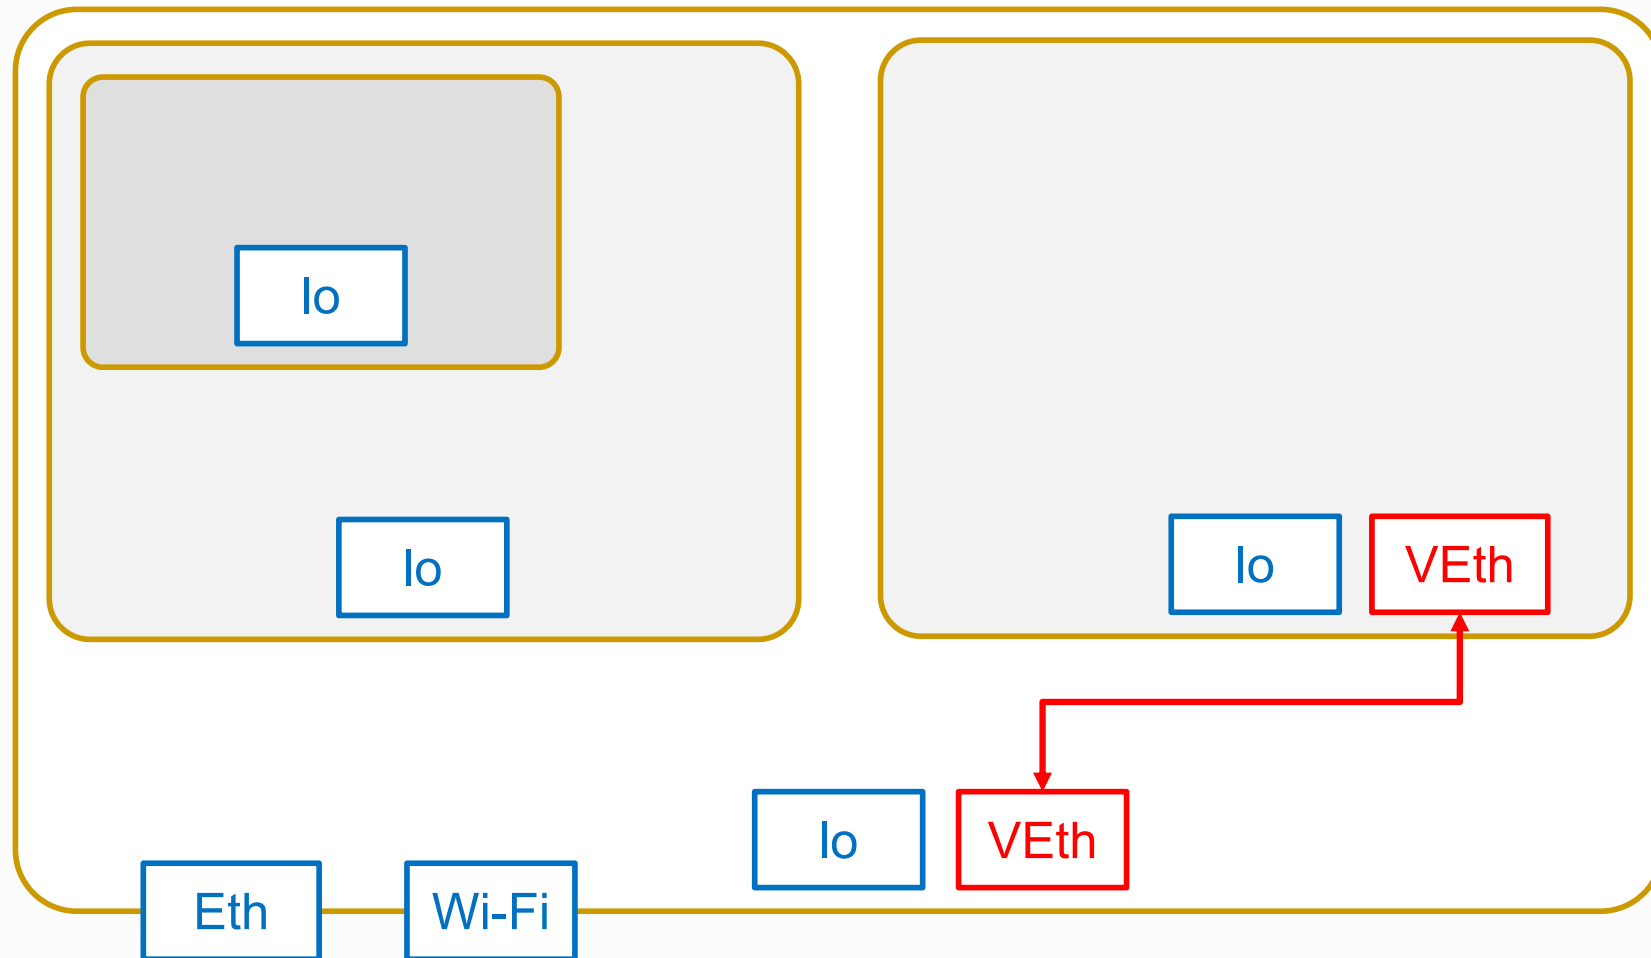
- ♦ **Network namespaces provide isolation of the system resources associated with networking:** network devices, Ipv* protocol stacks, IP routing tables, firewall rules, port numbers, etc;
- ♦ A physical network device can live in exactly one network namespace. When a network namespace is freed, its physical network devices are moved back to the initial network namespace;
- ♦ A **virtual network (veth) device pair provides a pipe-like abstraction** that can be used to create tunnels between network namespaces and can be used to create a bridge to a physical network device in another namespace. When a namespace is freed, the veth devices that it contains are destroyed.

Linux Programmer's Manual: *\$man network_namespaces*

Network namespace

- ▶ Linux network interfaces include
 - ◆ Physical interfaces (ETH, Wi-Fi, etc.)
 - ◆ Virtual interfaces (loopback, etc.)
- ▶ A network namespace is a separate set of network interfaces
 - ◆ All virtual
 - ◆ Including loopback

Network namespaces



Mount namespace

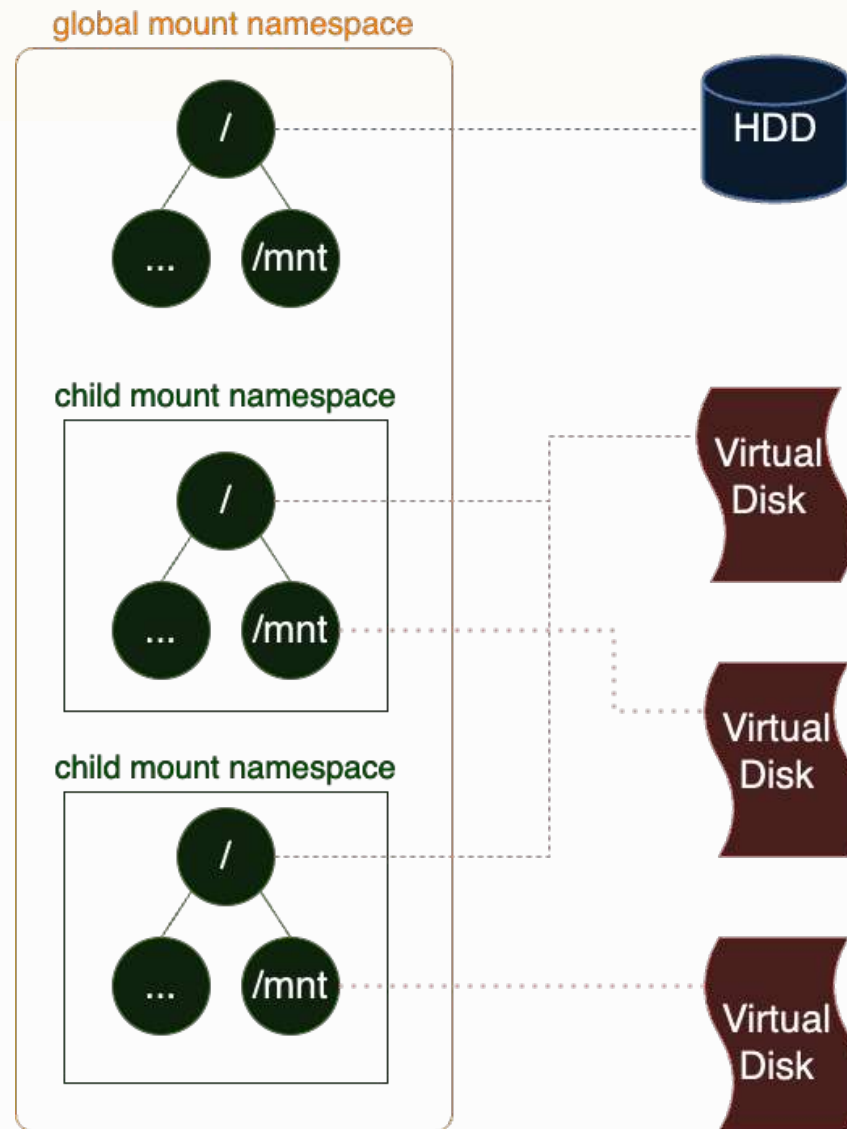
- ♦ **Mount namespaces provide isolation of the list of mount points seen by the processes in each namespace instance.** The processes in each of the mount namespace instances will see distinct single-directory hierarchies;
- ♦ The views provided by the (/proc/[pid])/mounts, mountinfo, and mountstats files correspond to the mount namespace in which the process with the PID [pid] resides;
- ♦ If a namespace is created using **clone**, the mount point list of the child's namespace is a **copy of the mount point list in the parent's namespace**. If it is created using **unshare**, the mount point list of the new namespace is a **copy of the mount point list in the caller's previous mount namespace**.

Linux Programmer's Manual: *\$man mount_namespaces*

Mount namespace

- ▶ A file system includes many mounts
 - ◆ Namely / is usually mounted on a storage device
- ▶ A new mount namespace allows processes in inner namespaces to change mount points without affecting outer namespaces
- ▶ This is beneficial to mount entire file systems to different devices or even files

Mount namespaces



UTS namespace

- ♦ **UTS namespaces provide isolation of two system identifiers: the hostname and the NIS domain name.** These identifiers are set using `sethostname` and `setdomainname`, and can be retrieved using `uname`, `gethostname`, and `getdomainname`;
- ♦ Changes made to these identifiers are visible to all other processes in the same UTS namespace, but are not visible to processes in other UTS namespaces;
- ♦ When a process creates a new UTS namespace using `clone` or `unshare` with the `CLONE_NEWUTS` flag, the hostname and domain of the new UTS namespace are copied from the corresponding values in the caller's UTS namespace.

Linux Programmer's Manual: `$man uts_namespaces`

UTS namespace

- ▶ Hosts have a name and a domain name
- ▶ A new UTS namespace allows processes on it to change those names without affecting the names in outer spaces
 - ◆ This is beneficial to “simulate” the execution of applications in arbitrary hosts

User namespace

- ♦ User namespaces isolate security-related identifiers and attributes, in particular, user IDs and group IDs, the root directory, keys and capabilities;
- ♦ A process's user and group IDs can be different inside and outside a user namespace. In particular, a process can have a normal unprivileged user ID outside a user namespace while at the same time having a user ID of 0 inside the namespace;
- ♦ In other words, the process has full privileges for operations inside the user namespace, but is unprivileged for operations outside the namespace.

Linux Programmer's Manual: *\$man user_namespaces*

User namespace

- ▶ Allows to create a process with all capabilities but with an arbitrary UID & GID mapping
 - ◆ The process keeps the UID and GIDs
 - ◆ But these need to be mapped to specific values
 - ◆ No mapping → 65534 (nobody)
- ▶ The mapping is a per-process, one time operation
 - ◆ `/proc/[PID]/uid_mapping`