



Minimal Information Disclosure with Efficiently Verifiable Credentials

David Bauer
Georgia Institute of
Technology
School of ECE
gte810u@mail.gatech.edu

Douglas M. Blough
Georgia Institute of
Technology
School of ECE
dblough@ece.gatech.edu

David Cash
Georgia Institute of
Technology
College of Computing
cdc@cc.gatech.edu

ABSTRACT

Public-key based certificates provide a standard way to prove one's identity, as attested by some certificate authority (CA). However, plain certificates provide a binary identification: either the whole identity of the subject is known, or nothing is known. We propose using a Merkle hash tree structure, whereby it is possible for a single certificate to contain many separate claims or attributes, each of which may be proved independently, without revealing the others. Additionally, we demonstrate how trees from multiple sources can be combined together by modifying the tree structure. This allows claims by different authorities, such as an employer or professional organization, to be combined under a single certificate, without the CA needing to know (or to verify) all of the claims. In addition to describing the hash tree structure and protocols for constructing and verifying our proposed credential, we formally prove that it provides unforgeability and privacy and we present performance results demonstrating its efficiency. As services move from user names and passwords to attribute-based identity verification, efficiency and scalability of claims verification will become a major issue. We have implemented a prototype client-server system, deployed the prototype in Emulab, and evaluated the server-side throughput for attribute-based identity verification. The results show that our approach can perform about 200 identity verifications per second, while the best competing approach can perform only about 2–5 verifications per second. Our approach is, therefore, better suited to today's high-volume Web-based services that demand the highest possible throughput.

Categories and Subject Descriptors

K.6.5 [Security and Protection]: Authentication

General Terms

Algorithms, Performance, Design, Security, Verification

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

DIM'08, October 31, 2008, Fairfax, Virginia, USA.

Copyright 2008 ACM 978-1-60558-294-8/08/10 ...\$5.00.

Keywords

Identity management, identity assertion, credential, hash-tree, Merkle tree, PKI, privacy

1. INTRODUCTION

Personal information is increasingly used to establish identity and authorize transactions in the digital world. At the same time, identity theft and fraud, based on unauthorized disclosure and misuse of personal information, are rampant [15], and individuals are increasingly concerned about providing this information to every digital entity with which they establish a relationship. The research described in this paper is based on several key principles of identity management. First, users should have the maximum control possible over what personal information of theirs is disclosed in any given interaction. Second, more reliance can be placed on personal information that is verified by trusted third parties than in self-reported information. And third, if verified personal information is to be used, mechanisms to prevent that information from being copied and misused by unauthorized parties are essential.

We assume an architecture in which there are identity providers that verify users' personal information and supply long-lived, off-line credentials that the users can show to service providers (also known as relying parties) [18]. Credentials are a common mechanism for verifying personal information in everyday life. Most people carry multiple physical credentials with them, from drivers' licenses to insurance cards to credit cards. A credential describes some set of attributes about the holder. For the obvious example, a driver's license states that the holder is licensed to drive a vehicle in the licensing state. However, due to drivers' licenses being almost universal, they are used as a general credential. As such, driver's licenses often include unnecessary information, such as the holder's date of birth, address, height, organ-donor status, and sometimes even social-security number. Electronic credentials can be simple, like a user-name and password, or more complex, like a public key infrastructure (PKI) certificate. A PKI certificate is an electronic document that holds an identity and a public key, and that is signed by a certificate authority (called the issuer). The user holds the associated private key to prove that they are the legitimate holder of the certificate. The user-name and password combination is the most widely used scheme, because of its simplicity. However, this scheme also provides no direct information about the user. A user-name must be attached to a previously made account, and some other form of credential must be used to tie an identity

to the account. Such accounts are very seldom shared between different domains, leading users to accumulate many different accounts, often with different user-names and passwords. PKI certificates for users are less common, but can solve the problem of needing to keep track of many different user-names and passwords.

Minimal information disclosure in any given interaction is desirable from a user's perspective, and may even be necessary for a given technology to be widely adopted [9]. Clearly, if a user wants to release the minimal amount of personal information on a given interaction, this rules out a single credential approach, in which each user maintains one credential containing all of their personal information and uses that credential for every interaction. As in [9], we say that a user makes a claim about herself when she gives information about one or more personal attributes to a digital entity. Due to the wide variety of personal information that is used in digital interactions, the number of different possible claims is extremely large. The problem to be solved is, therefore, to provide an efficient and reliable mechanism that allows users to assert arbitrarily many (or at least a large number of) verifiable claims over a sequence of interactions with different digital entities.

An example of minimal information disclosure is the problem of verifying that a user is at least 18 years of age. Clearly, verifying the user's date of birth is sufficient but not necessary, and would reveal a very sensitive piece of personal information that could assist an identity thief in masquerading as that user. We refer to the claim that a user "is at least 18 years of age" as a micro-claim, with the (macro-)claim in this case being "the user's date of birth is xx/xx/xxxx". Many different micro-claims can be derived from a single claim, e.g. the user is "at least 18", "at least 21", "at least 35", "at least 65", etc. Another example of the macro- to micro- expansion would be a street address, which can be broken into country, state, county, city, zip code, neighborhood, street name, street number, apartment number, congressional district, school district, etc. In a more technical example, the micro-claims could consist of a large number of one-time use tokens, such as one-time use credit card numbers.

We adopt an approach where a large set of micro-claims is enumerated statically, and updated dynamically as needed. Instead of generating and maintaining a single credential per micro-claim, which is extremely inefficient from the standpoints of storage space, bandwidth, and computation time, we propose a method wherein a single credential can be maintained that allows the user to dynamically specify an arbitrary subset of micro-claims for a given interaction without revealing the other micro-claims. We expect well-defined standards and automated tools to handle most of the details associated with the micro-claims. We also include in this credential, and the protocols that use it, mechanisms to make it difficult for an attacker to make a copy of the credential and masquerade as the user. In addition to allowing the user to update her set of microclaims as certified by one authority, the credential allows information that is verified by different identity providers to be combined in one structure. This allows users to spread out their personal information across different identity providers, thereby lowering their risk when one of their identity providers experiences a security breach. The details of this new credential mechanism and its associated protocols are provided in the remainder of this paper.

Defining the level of privacy needed by users is a subtle and hotly debated issue. One attempt to quantify privacy is the idea of linkability. Two things—events, transactions, credentials, claims, or users—are linkable if they are known to have an underlying connection. It is necessary for some things to be linked; for example, if a claim is not linked to any user or evidence, it is meaningless. But more often the focus is on unlinkability—when things can't be recognized or shown to be connected, even when they are. For example, consider buying an item with cash versus buying an item with a credit card. When using cash, there is usually no record of who bought the item. If several items are bought in separate transactions, it usually cannot be determined that they were bought by the same person. When using a credit card, there is a clear record of who bought the item. Items purchased at different times and even at different places may be identified as having been bought by the same person. Additionally, the real name/identity of that person is known and usually printed right on the receipt. This is a clear invasion of privacy, and yet credit cards are still extremely popular. Credit cards are popular because they are convenient, and this invasion of privacy is accepted because it is necessary to reduce fraud and abuse in the system.

Our goal here is to design systems that preserve the user's privacy as much as possible, while remaining auditable from the viewpoint of the service providers. That is, a user should have control over which claims are presented, but we do not attempt to make repeated uses of the user's credential unlinkable. This concession to social practicality gives a strong engineering advantage compared to the other credential systems discussed in the related work section: we are trying to solve a somewhat easier problem. As such, we can focus on making the system faster and more efficient. While complete anonymity, e.g. in the form of unlinkable transactions, is a noble goal, it is our belief that many types of service providers will not accept it as a practical solution. Our approach, by contrast, tries to improve on the current practice by allowing users to minimize the amount of personal information they entrust to service providers while recognizing that release of some personal information is inevitable.

The disclosure of personal information carries with it the risk that users may be tracked across multiple domains via an external correlation of their data (for example, a mailing address). Such a correlation attack depends only on the data released and not how the data is released, and so we do not consider it further.

2. BACKGROUND

2.1 Expected Scenario

We consider a scenario with three types of parties: identity providers, service providers, and users. An identity provider is an entity in a position to make authoritative statements about a user. An identity provider can be a third party certificate authority, a government office, an employer, a professional organization, etc. A service provider is an entity that wants or needs to check the identity of users. A service provider can be any organization, government office, or individual. The user is the holder of the credential, and is always a person in our scenario.

We acknowledge that it is notoriously difficult to deploy a usable and secure PKI to bind keys to users, but it seems

necessary for off-line credentials. The functionality provided by verified claims should serve as intuitive and strong motivation for users to choose these methods over current practice.

Revocation is an issue for any off-line credential system. Possible solutions include periodically distributed revocation lists, live checking of a credential's status, and expiration of credentials in a timely manner. All of these are higher-level design concerns, and do not effect the underlying credential system as described here.

This paper only considers off-line credentials, where the identity provider is *not* involved with every showing. Compared with on-line credentials, off-line credentials are advantageous in situations where the user and relying party are not connected to the Internet, where the user does not want to contact the identity provider for each transaction for privacy reasons, or where it is not economical for identity providers to maintain 24/7 availability.

2.2 Requirements

The design of our electronic credential was driven by several requirements. First, the credential must not simply hold a single identifier. Identity is a complex matter, and often a name or serial number is not what is important in one's identity. Second, the user of the credential—the one whose identity is being proven—must have as much control over the process as possible. Considering the first requirement, the user should be able to select which attributes of their identity are released to a particular entity. Third, the credential must have some form of copy protection. Using the credential should not place it at risk of being copied. Fourth, neither the user nor the service provider should have to contact the identity provider to verify the credential. Finally, the credential must be memory and computationally efficient to store and use.

2.3 Related Work

Digital credentials have been proposed before in various forms. Some are pseudonym systems whereby a user can present a different pseudonym to each service provider. The pseudonyms are linked cryptographically such that even colluding service providers cannot link together two pseudonyms belonging to a user, but the user can use a credential assigned to one pseudonym with another pseudonym. Some digital credentials are claims based, whereby the credential contains assertions about the user, but no pseudonym. Our proposed credential is claims based.

David Chaum is credited with proposing the first digital credential system [11]. Chaum's system is a pseudonym system, or more precisely the first description/definition of a pseudonym system. Chaum's original proposal did not describe how to implement such a system. A method for implementing the system was later described by Chaum and Evertse based on the RSA crypto-system and a semi-trusted third party [12].

Stefan Brands developed a different form of digital credential, in which a user has a single public credential, but that credential is pseudo-anonymous, even to the issuer [3]. The credential holds attributes that the user can selectively prove to a service provider. Repeated showings of the same credential are linkable, however. Since the credential is issued blind by the identity provider, effectively a user has one global pseudonym. The credential can be reissued easily, allowing

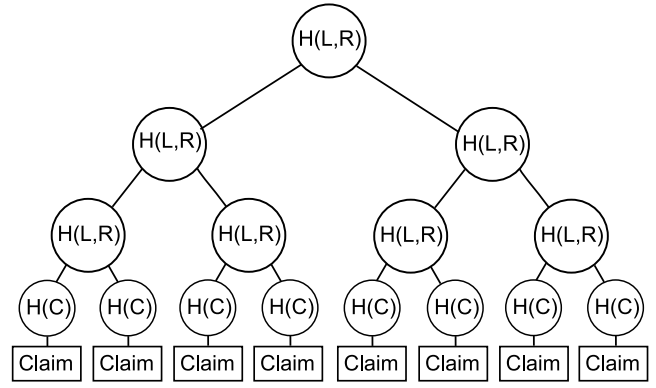


Figure 1: Merkle hash tree with leaf nodes holding hashes of claims.

the user to change the global pseudonym, as permitted by the identity provider. Brands emphasizes that showing a credential is done by zero-knowledge proof, but that has no impact on the comparisons in this paper. Credentica's U-Prove Software Development Kit is based on Brands' work [4].

Camenisch, et al., have proposed and implemented yet another form of digital credential, or more precisely, yet more forms [6, 7, 8]. While they describe a system for implementing a Chaum-like pseudonym system, their system is much more flexible, and can be used without pseudonyms. While having stronger anonymity properties, the algorithms are also slower than Brands' credentials according to [5]. IBM's idemix system is based on Camenisch, et. al.'s work [6].

While their aim is very different, the closest design to our credential system is the redactable signature scheme described by Johnson et al in their paper on homomorphic signature schemes [16]. A comparison between our credential and their signature scheme is given in Section 3.4.

3. DESIGN

3.1 Merkle Hash Trees

The proposed credential is based on Merkle hash trees [17] and standard public-key infrastructure (PKI) certificates. A Merkle hash tree is essentially a binary tree where each internal node holds the hash of the concatenated values of its two children nodes. The leaf nodes hold the data of interest. In this way, a large number of separate data can be tied to a single hash value. In addition, by storing the internal node values (or a subset thereof), it is possible to verify that any of the leaf nodes is part of the tree without revealing any of the other data. Ralph Merkle first introduced this structure as a way to efficiently handle a large number of Lamport one-time signatures. It has since been adapted for uses such as the large-scale time-stamping of documents [2] and tracking data in peer-to-peer networks [10]. A basic Merkle tree is shown in Figure 1.

3.2 Credential overview

The credential consists of two parts: a public part and a private part. The public part of the credential is a certificate. The certificate holds information about the issuer, the certification chain for the issuer, the type of certificate, the date range over which the certificate is valid, the user's public

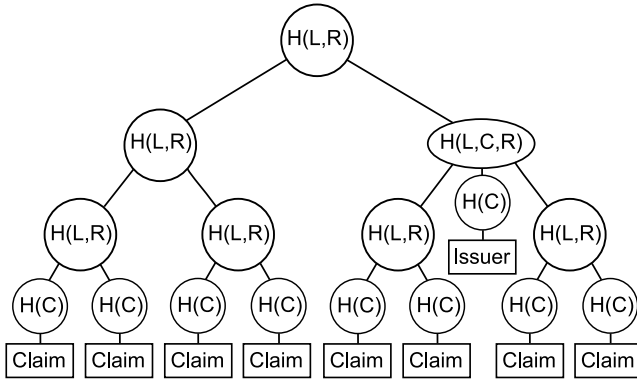


Figure 2: Modified Merkle hash tree with subtree.

key, and the root hash of a Merkle hash tree. The certificate should not hold any data about the user directly, even data as common as a name. As per standard operation, the certificate will be signed by some certificate authority, which is an identity provider in this system. The private part of the credential consists of a private key and a Merkle hash tree whereby all of the leaf nodes are (hashes of) micro-claims about the identity of the user, who is the credential holder. The Merkle tree structure allows the credential holder to prove any subset of the claims in the tree, with only the single signature on the certificate.

As an additional improvement, a slight modification to the structure will allow the use of a single credential containing claims from a variety of identity providers, without requiring one identity provider to verify all of the claims. Consider again the Merkle tree of claims, but with subtrees coming from different identity providers. For example, one subtree could contain claims certified by a government registry, while another subtree could contain claims certified by an employer. In the basic structure, the identity provider must either see all of the claims or trust the providers of all of the subtrees that they only contain claims relevant to their topic. Neither solution is ideal. To provide a third option, consider adding an optional third branch to some internal nodes of the full tree. (These nodes correspond to root nodes of the subtrees.) The third branch contains a certificate for all of the claims in the subtree rooted in the parent node. When verifying a claim which is in such a subtree, the form is different—three hashes are concatenated in a node instead of two—so the verifier knows that this claim is under a different certificate than the credential as a whole. For example, in Figure 2 the left hand subtree contains claims certified by the overall identity provider. The overall identity provider is referred to as the certificate authority (CA). The right hand subtree contains claims certified by some other party, which do not have to be verified by the top level CA. From here on, the term “subtree” is used to refer to a branch containing nodes from a different identity provider, rooted in one of these sub-roots.

We would like to identify a node as being a leaf node without seeing the associated claim. This is required for the security proof in the appendix. This discrimination is achieved by simply appending a bit to the end of the hash, either a one if the node is an inner node, or a zero if the node is a leaf node. For efficiency, our implementation simply overwrites the least significant bit of the hash value, instead of append-

ing another bit. This is equivalent to truncating the output of the hash function used by a single bit (with a insignificant decrease in security), and then appending the indicator bit.

In order to keep the system secure while still only needing to check the signature on a single certificate, nested subtrees are forbidden. If a service provider encounters a claim in a nested subtree, it should reject the credential.

A tree might have no claims outside of subtrees. In this case, the top level CA is signing that it has verified that the form of tree is correct, that all certificates for subtrees are valid, and that the user possess the private key(s) matching the subtree certificates and the top level certificate. The CA could be fully automated in this situation, allowing users to easily update their credentials when subtrees are added or modified. Note also that changes to attributes in one subtree only require recomputation of that subtree, the main root hash, and the top-level signature. Recomputation of other sub-trees and other signatures is *not* required.

3.3 Protocols for the Credential

Creating a credential is both conceptually and computationally easy. In the case where there is a single identity provider for the credential, there are roughly four steps required: first, agree on a list of claims. Next, generate the hash tree for the claims. Third, verify that the user possesses the private key. Finally, produce and sign the public certificate.

First, the user and identity provider agree on a list of claims. The logistics involved in an identity provider verifying the user claims are beyond the scope of this paper. Second, either the user or the identity provider generates the hash tree for the claims. Random padding must be added to the claims before they are hashed, as discussed shortly. With a single identity provider, the tree will always be balanced, bounding the number of interior nodes to the number of claims. The number of hashes needed to generate the tree is therefore bounded to twice the number of claims—generally an insignificant amount of computation time. Third, the identity provider must verify that the user holds the private key matching the public key of the credential. The user can reuse an already generated key pair or generate a new key pair for the credential. The identity provider does not need to ever know the private key. Finally, the identity provider creates and signs the public certificate for the credential.

Random padding must be added to the claims before they are hashed in order to prevent dictionary attacks against parts of the tree. The random padding can be generated and stored in a number of different ways. For example, Johnson, et al. [16], describe a method credited to Goldreich, Goldwasser, and Micali [14] that uses a pseudo-random function whose output is twice the length of its input. A single seed value at the root node is expanded as per the tree branching structure using the pseudo-random function. Each node therefore has a pseudo-random value associated with it that can be used to generate the pseudo-random values for all of its children nodes. A computationally easier method is to use a seed value and the claim index with a pseudo-random function to generate the padding for each claim. Neither of these methods will work well for our system, unfortunately, due to the inclusion of different subtrees within a larger tree. In the most general case, simply storing the random padding for each claim is probably the easiest solution. The overhead is small, just 10-16 bytes per claim. Alternative approaches

are possible, but all require the user to have a global secret, keep the private keys for all certificates, or store extra private information with each credential.

Creating a credential with claims from multiple identity providers is done by integrating credentials as subtrees. One identity provider will be the final one, referred to as the certificate authority (CA). The user creates credentials with all of the identity providers, except for the final one, by the procedure described above. Then, the user creates a credential with the CA. For this final credential, the hash trees from the other credentials are incorporated into the final hash tree. The root nodes for the incorporated subtrees will have a third branch added, each containing the certificate from the original identity provider of that subtree, as shown in Figure 2. (The root hash values in the subtree certificates will no longer match, but that doesn't matter.) All subtrees in the credentials being added must be removed from their trees and added separately to the top level tree. This prevents nested subtrees. The CA verifies the structure of the tree, including the top two leaf nodes of each subtree (as they are needed to calculate the sub-root's hash), as well as the public part of the sub-credentials, and the associated public/private key pairs. In general, the public/private key pairs may be the same as or different from each other and the top level key pair. All that matters is that the user holds the corresponding private keys at the time that the credential is assembled. In a more restricted setting, the public/private key pair may be required to be the same for all subtrees. The CA does not need to see the claims from subtrees.

To provide an example of combining credentials, imagine a user, Alice, who already has a credential containing claims from three different identity providers. Alice is then issued a credential from her employer. The two credentials are shown in Figure 3. Alice's employee credential already has a number of subtrees, containing separate credentials from the engineering department, human resources department, and her own lab. The credential also has a top level claim of her employee ID. Alice goes to a certificate authority to create a new credential from these two credentials. She submits no claims to be verified by the CA and seven sub-credentials to be included in the final credential—three of these are from her previous credential, three were embedded in her employee credential, and the last is the credential claiming her employee ID. The CA will verify the seven certificates, verify that Alice possesses the appropriate private key(s), generate a new tree, and provide Alice with the new certificate.

The protocol for using the credential follows conventional PKI certificate usage. The user connects to a service provider, either over a wide or local area network, and requests some service, sending the public part of the credential. The service provider requests the appropriate identity attributes from the user. The user provides the claims that match the requested attributes and the intermediate node values and path information necessary for the service provider to verify each of the claims. If any of the claims are in a subtree certified by a different identity provider, the accompanying certificates must be included with the set of claims. The service provider verifies that the claims are in the hash tree specified (via the root hash) in the certificate part of the credential. The service provider also verifies the signature on the certificate part of the credential. In order to verify that

the user is the holder of the credential, the service provider also verifies that the user possesses the private key which matches the public key claimed by the credential. This can be done by standard methods, such as challenge/response or as part of a secret key agreement operation, as long as the key verification is tied to the specific claims being asserted by the user. This can be done, for example, by hashing all of the claims asserted by the user as part of the challenge field. In addition to the cryptographic verifications, the service provider must of course confirm that it trusts the identity providers to assert the claims in the credential. For example, a claim of an individual's address asserted by the Bar Association would be out of place. Similarly, an assertion by the Department of Motor Vehicles that an individual was a licensed lawyer should not be trusted.

The procedures for creating and using credentials given herein are purposely generic. We prefer to constrain the structure and properties of the credential and leave some flexibility in the procedures for generating and verifying it. As a prototype, we have developed specific implementations of these procedures, and performance results based on these implementations are reported in Section 5.

3.4 Comparison to Related Systems

Johnson, et al., define and construct redactable signatures using a Merkle hash tree to allow the signature verification of a message even when parts of the message have been deleted [16]. In their example, a body of text is signed such that it can be redacted at some level of granularity (sentence, word, or character making the most sense). They do not appear to have considered using the construct in the context of a credential. While both systems share a core idea—using a Merkle hash tree to hide signed elements—the two systems have significant differences. Both systems use slightly different hash functions (constructs) for the leaf and inner nodes. In their signature scheme, the input to the hash function is specified to match a certain form, which is important to their proof of security. In our credential system, the output of the hash function is specified to match a certain form, which is sufficient to meet the requirements of the security proof, while providing the property of being able to identify a leaf node from its hash value alone. Subtrees of the type used in the credential system are not used at all in the signature scheme. To the best of our knowledge, the idea of combining trees by modifying the tree structure as we do is novel, and introduces extra complications. The three-child inner nodes (shown in Figure 2) and modified hashing scheme (discussed in Section 3.2) both result from the idea of combining subtrees from different identity providers. Combining subtrees allows claims from many different sources to be in a single credential, under a single certificate.

Brands mentions in his book the idea of using hash trees to store claims, but dismisses it because it does not provide the properties he desires [3].

3.5 Copy Resistance

The credential is strongly copy resistant under normal operations, due to the incorporated public/private key pair. Whenever the credential is used, the private key is used to prove that the user is the authorized credential holder. Since the private key never leaves the user's machine, it can't be copied by the service provider or any third party listening to the exchange. (The private key can of course be copied if




Figure 3: Combining a generic credential with an employee credential.

the user’s machine is compromised, but that is beyond the scope of this paper.)

This copy resistance property extends to preventing standard man-in-the-middle attacks. For example, consider how a phishing site handles conventional one-time passwords or other two-factor authentication methods. A user logs in to what they believe is the legitimate site of their bank, broker, or other service provider, but which is really a phishing site. The user enters their user-name and password and then, either in the same step or in a second step, enters a one-time password from a sheet of paper or electronic device. The phishing site simply passes all the information from the user on to the site being spoofed, and returns the responses from the legitimate site to the user. After the user has successfully logged into the legitimate site, through the phishing site, then the phishing site is logged in as the user and can perform whatever actions it wants.

The standard protection against this type of man-in-the-middle attack is the use of server-side certificates with SSL/TLS. When the user initiates a secure connection to a service provider, the service provider replies with a PKI certificate. The user’s client (web browser) usually handles checking the validity of the certificate automatically, interrupting the user only when a problem is detected and simply showing a non-intrusive indicator when the secure connection is setup without problems. Phishing sites have used a variety of techniques to get around server-side certificates, including simply not using SSL/TLS, making the browser look like it is using SSL/TLS when it is not, and obtaining valid certificates for different domains, which can be mistaken for the legitimate domain by an unwary user. All of these tricks are possible because of the disconnect between the authentication system (SSL/TLS), and the thing being authenticated (the service provider). The certificates authenticate the domains used, instead of the actual service providers.

In contrast to the server-side certificates, our credentialing scheme intimately ties the authentication system (the private/public key pair and the certificate as a whole) to what is being authenticated (the credential as a whole). Consider again the phishing site, but with the credential in place of the user-name, password, and one-time password. When the user connects to the phishing site, the phishing site can relay requests to the legitimate service provider and replies back to the user as before. However, when the service provider sets up a secure connection, it will use the user’s certificate to setup the session. The phishing site can substitute its own certificate, but then it is no longer impersonating the user. Alternatively, the phishing site can pass on the user’s certificate, but then the phishing site loses control of the

session. The service provider can setup an authenticated tunnel to the user, so that the phishing site can no longer modify the traffic without being detected.

4. SECURITY

In this section we discuss some potential attacks against implementations. We provide a formal analysis of the security properties of the underlying cryptographic construction in the appendix.

4.1 General Security Discussion

The basic attacks that the credential must resist are forgery, theft, privacy compromise, man-in-the-middle, and collusion.

Forgery - a malicious user should not be able to forge a valid credential containing invalid claims.

Theft - the credential should be resistant to theft.

Privacy compromise - an attacker should not be able to learn more about the user of a credential than the user chooses to reveal.

Man-in-the-middle - a hostile service provider should not be able to misuse or reuse a user’s credential that is presented to it.

Collusion - two or more users working together should not be able to make a claim that no single user in the group can make alone.

During the design process, several specific attacks against earlier versions of our system were discovered. They include dictionary attacks against neighbor nodes, combined replay/man-in-the-middle, hidden subtrees, and the broken hash attack.

Dictionary attacks against neighbor nodes come from the fact that when a credential is used, the hash values of unreleased claims must be provided. Assuming a secure hash function is used, an attacker cannot determine the value of the unreleased claim via cryptanalytic attack. However, since most claims are likely to be in a standard form, a dictionary attack should often be successful. We protect against this type of attack by padding the claims with random or pseudo-random data.

A combined replay/man-in-the-middle attack may be possible when the verification that a user holds the private key for a credential is completely unrelated to the showing of the tree for the credential. In this situation, an attacker can perform a replay attack against the showing of the tree, and then a man-in-the-middle attack against the user proving possession of the private key. While this attack either does not apply or would be meaningless to many uses of the credential, it is still a possible attack. We protect against

this type of attack by requiring that the verification of the private key be linked to the specific claims being shown.

To maximize the efficiency of showing a credential, we would like to minimize the number of slow, public-key operations performed. The best that we can do is to have only a single certificate verification and a single public/private key verification, regardless of the number of claims and how many different subtrees the claims may be in. During the issuance of a credential an identity provider will check the certificates for all subtrees within the tree. Therefore, when the credential is shown, the service provider does not need to recheck those certificates, as long as it trusts the top level CA. However, consider the case of a nested subtree. An identity provider shouldn't be able to see the claims in subtrees, due to privacy requirements. But this means that a subtree can hide another subtree. In particular, a subtree issued by a bad (but untrusted) identity provider could hide a subtree apparently issued by a trusted identity provider, but in actuality was forged and has a bad signature. This could be prevented by several different protocol changes. The most efficient of these is by preventing the inclusion of a hidden subtree in the first place. Our current recommendation is simply to ban nested subtrees altogether, and have the service providers ensure that credentials which have nested subtrees are rejected.

Another structural attack that was encountered is implementation specific. In our prototype implementation, plain X.509 certificates are used. X.509 certificates are designed to match a directory entry for a particular subject, and their structure is fairly rigid around that purpose. Therefore, our implementation didn't specify any extra details about the hash tree—such as the hash algorithm used—in the certificate itself, but instead stored that information with the tree. This can lead to an attack if a hash algorithm trusted for use in creating the tree is broken. Given the right conditions, an attack could create a tree using the broken algorithm, which could then collide against the root hash stored in a valid certificate, even if the valid certificate was for a tree using a different (and hopefully more secure) hash algorithm. We consider this attack to be fairly minor, because it requires a primitive of the system to be broken badly—in order to find a collision from one hash function to another should effectively require a preimage attack, and not a simple collision. As version 3 of X.509 certificates allows arbitrary extensions, the hash algorithm can be specified in the certificate, anyway.

5. PERFORMANCE

Our minimal disclosure credential design focuses on conventional operations (one-way hashes) and minimizes the number of asymmetric/public key operations. This allows our system to be much faster than the digital credential schemes by Brands and Camenisch, et al., which require many more of these expensive public key operations. In this section, we report on two sets of performance experiments that validate this claim. In the first set, we evaluated the computation time to verify credentials of different types with a varying number of claims. This set focuses solely on the computational burdens placed on servers, which is likely to be the most significant performance bottleneck in a system. In the second set of experiments, we implemented and deployed an actual client-server system in Emulab [19], and evaluated the server throughput. This set of experiments in-

Operation	Time (μ -sec)
Verify Tree (SHA-256)	
1 claim of 2048	31
20 claims of 2048	100
All 2048 claims	8030
Verify Certificate	
RSA, 1536 bit	620
Verify Brands	
1 claim	38,000
20 claims	283,000
2048 claims	26,000,000

Table 1: Time efficiency of hash tree and certificate

cluded communication costs and additional client-side and server-side computations costs in addition to credential verification time, and it has multiple clients accessing the service simultaneously, to provide a realistic overall assessment of performance. In both sets of experiments, we compared our hash tree approach against Brands' credentials and the brute force solution of embedding each claim in a separate certificate. For Brands' credentials, we used the parameters from [5, Section 2.2] ($|p| = 1600$, $|q| = 256$, $|s| = 160$).

The first set of experiments (credential verification time only) was performed on an Intel Core 2 Duo E6600 running at 2.4 GHz using Sun's server JVM version 1.6.0. In all cases the operations were run repeatedly before timing, so as to force the JVM to fully optimize the operations. Our prototype uses X.509 certificates (with the root hash as the distinguished name) for the public key certificate part of the credential. The certificate is signed using RSA, and the CA's key is 1536 bits. Results are shown in Table 1. The total time to verify one of our hash tree credentials is equal to the tree verification time plus the time to verify one certificate. Trees containing claims from multiple authorities would require verifying the hash tree and one certificate for each authority that has at least one claim being asserted. The hash tree used in these experiments contains 2048 claims. From Table 1, we see that verifying one claim with the hash tree approach is approximately 58 times faster than verifying one claim in a Brands' credential. For 20 claims, the hash tree approach is approximately 390 times faster than Brands', and with 2048 claims, the speedup factor is approximately 3,000. Speedup compared to the brute force approach of one certificate per claim is also substantial. This speedup is linear with the number of claims, i.e. approximately 20 times faster with 20 claims, 100 times faster with 100 claims, etc.

Our second set of experiments was done using an actual client-server implementation deployed in Emulab [19]. In these experiments, there were a variable number of clients and one server, which accepted connections from the clients and verified one credential per client request. Each client and server instance was run on a separate 3 GHz Pentium Xeon. Each client generated verification requests one after another, waiting for one response and then immediately generating a new request. We again waited until the server had processed a number of requests before timing the operations to ensure full JVM optimization. All parameters were the same as in the first set of experiments. Here, our primary performance measure is the server throughput (number of verifications done per second). The results of these experiments are shown in Figure 4. The figure shows

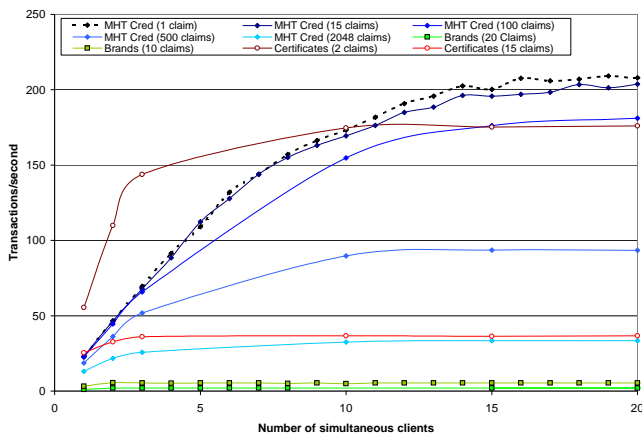


Figure 4: Credential Verification Throughput vs. Number of Clients

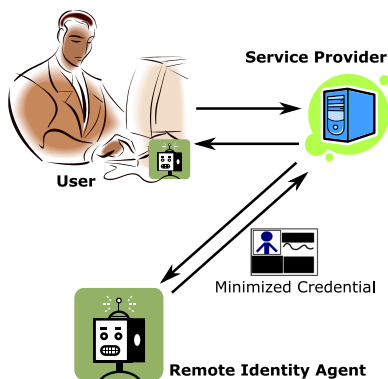


Figure 5: Remote identity agent system.

that throughput with Brands’ credentials saturates at only a few clients, while the hash tree credentials can handle up to about 15 clients before saturating. The peak throughput for the hash tree credentials is approximately 210 verifications per second, while the peak throughput with Brands’ credential having 10 claims is approximately 5 operations per second and with 20 claims it is only about 2 operations per second. Adding claims to the hash tree credential has only a modest impact on performance. With 15 claims, the peak throughput is still over 200 verifications per second and even with 2,048 claims, the throughput is still around 33 verifications per second. The figure also shows that throughput for the brute force approach with one certificate per claim drops off rapidly as the number of claims increases. Overall, the Emulab results support the earlier experiments in that the hash tree credentials retain a strong performance advantage relative to Brands’ (approximately 40-100 times faster for moderate numbers of claims).

6. USE OF CREDENTIALS IN A DEMONSTRATION SYSTEM

The prototype credential system is already being used by a second prototype system—one working with a remote identity agent. Figure 5 shows the current remote agent prototype design in use. The user is issued one or more credentials by an identity provider. The hash tree and claims of the cre-

entials are stored on a remote, networked attached identity agent. The public key or keys are stored on the user’s local device. The user proves ownership of the credential using the private key, while the bulk of the credential is sent from the remote agent to the service provider, after having received appropriate authorization from the user. As we are allowing for a resource-constrained user device, such as a cell phone, we try to limit the number of network connections by sending the authorization for the remote agent through the service provider. The authorization is encrypted and signed to prevent tampering and protect the user’s privacy, of course. The service provider validates the credential as usual.

The system is implemented in Java. The user’s local device is a local agent program spawned from a webbrowser, while the service provider is a modified webserver. A video demonstrating its use is available on our website.[1]

7. CONCLUSION

Privacy is important both as a protective principle and as a security measure. Identity theft is a serious and widespread crime. The Federal Trade Commission reports that over a quarter of a million identity theft complaints were received in 2005, in addition to over 430,000 other fraud complaints. Internet-related complaints accounted for almost half of those [13]. Protecting personal information is vital to reducing identity theft. Limiting information disclosure does not require that accesses to service providers be completely anonymized. A user may repeatedly present the same credential or service provider specific claims to a service provider. This could facilitate many useful applications such as user-controlled services for monitoring usage of their credentials. In addition, for security sensitive applications and with appropriate safeguards in place, accumulated user information could be sent to the identity providers to enable auditing of credential use.

8. ACKNOWLEDGEMENTS

The authors wish to thank Alexandra Boldyreva for helpful discussions, and Ralph Merkle for inspiration. This research was supported in part by the National Science Foundation (under Grants CCR-0208655 and CNS-CT-0716252), Nortel Networks, the state of Georgia, and the Institute for Information Infrastructure Protection. This material is based in part upon work supported by the U.S. Department of Homeland Security under Grant Award Number 2006-CS-001-000001, under the auspices of the Institute for Information Infrastructure Protection (I3P) research program. The I3P is managed by Dartmouth College. The views and conclusions contained in this document are those of the authors and should not be interpreted as necessarily representing the official policies, either expressed or implied, of any of the sponsors.

9. REFERENCES

- [1] D. Bauer. Video demonstration of the credential-holding remote identity agent. http://users.ece.gatech.edu/~gte810u/RIIdA_Video/, 2007.
- [2] D. Bayer, S. Haber, and W. Stornetta. Improving the efficiency and reliability of digital time-stamping. In *Sequences II: Methods in Communication, Security*,

and Computer Science, pages 329–334. Springer-Verlag, 1993.

- [3] S. Brands. *Rethinking Public Key Infrastructures and Digital Certificates; Building in Privacy*. MIT Press, 2000.
- [4] S. Brands. Credentica - u-prove sdk, 2007.
- [5] S. Brands, L. Demuyne, and B. D. Decker. A practical system for globally revoking the unlinkable pseudonyms of unknown users. In *(Accepted to) 12th Australasian Conference on Information Security and Privacy*, 2007.
- [6] J. Camenisch and E. V. Herreweghen. Design and implementation of the idemix anonymous credential system. In *Proceedings of the 9th ACM conference on Computer and communications security*. ACM PRESS, 2002.
- [7] J. Camenisch, S. Hohenberger, and A. Lysyanskaya. Compact e-cash. In R. Cramer, editor, *Advances in Cryptology — EUROCRYPT ’05*, volume 3494 of *Lecture Notes in Computer Science*, pages 302–321, 2005.
- [8] J. Camenisch and A. Lysyanskaya. An efficient system for non-transferable anonymous credentials with optional anonymity revocation. In *Proceedings of the International Conference on the Theory and Application of Cryptographic Techniques: Advances in Cryptology*. Springer-Verlag, 2001.
- [9] K. Cameron. The laws of identity, 2005.
- [10] J. Cates. Robust and efficient data management for a distributed hash table. Master’s thesis, Massachusetts Institute of Technology, May 2003.
- [11] D. Chaum. Security without identification: transaction systems to make big brother obsolete. *Communications of the ACM*, 28, Oct 1985.
- [12] D. Chaum and J.-H. Evertse. A secure and privacy-protecting protocol for transmitting personal information between organizations. In *Advances in Cryptology*, pages 118–167. Springer-Verlag, 1987.
- [13] Consumer fraud and identity theft complaint data, 2006.
- [14] O. Goldreich, S. Goldwasser, and S. Micali. How to construct random functions. *J. ACM*, 33:792–807, 1986.
- [15] D. Hardt. Identity 2.0, 2005.
- [16] R. Johnson, D. Molnar, D. X. Song, and D. Wagner. Homomorphic signature schemes. In *Topics in Cryptology – CTRSA 2002*, volume 2271, pages 244–262. Springer-Verlag, 2002.
- [17] R. Merkle. A certified digital signature. In *Advances in Cryptography*, pages 218–238. Springer-Verlag, 1989.
- [18] Microsoft. Microsoft’s vision for an identity metasytem, 2005.
- [19] B. White, J. Lepreau, L. Stoller, R. Ricci, S. Guruprasad, M. Newbold, M. Hibler, C. Barb, and A. Joglekar. An integrated experimental environment for distributed systems and networks. In *Proc. of the Fifth Symposium on Operating Systems Design and Implementation*, pages 255–270, Boston, MA, Dec. 2002. USENIX Association.

APPENDIX

A. FORMAL ANALYSIS OF SECURITY

We provide a rigorous analysis of our credential system. We first formally define the new notions of *unforgeability* and *privacy* for a credential system and show that our system achieves these under standard cryptographic assumptions.

Unforgeability states that a user cannot convince a service provider that a set of claims is true unless those claims were actually approved by an identity provider. We formalize this in a strong way: we allow adversaries to adaptively obtain an unbounded polynomial number of credentials on sets of claims, and then require that the adversary can not generate a credential on a set that was not contained in one of the valid credentials. Furthermore, the adversary may request that credential trees be combined. The adversary is also considered to have won if it can fool a service provider into accepting a valid claim from one identity provider as a claim from another identity provider which did not approve the claim. For simplicity, we analyze the cases where the claim tree contains no subtrees or is composed entirely of subtrees. The analysis can easily be extended to the case where claims are mixed with subtrees.

Privacy states that no partial information about unrevealed claims is leaked. We will allow adversaries to adaptively obtain credentials and then submit two sets of “challenge” claims of equal size. One of the challenge claim sets is selected at random and then a credential is generated on that set and returned to the adversary. The adversary can ask that parts of the challenge credential be revealed before outputting a guess. Of course, we must restrict the adversary to revealing only claims on which the sets match (otherwise distinguishing is trivial).

We leverage two properties in our analysis that were not considered in the redactable signature scheme of Johnson, et al., which used a similar technique of constructing a Merkle tree and then signing the root [16].

First, we consider the privacy of unrevealed claims, a notion that did not apply to the context of redactable signatures. If we translate Johnson, et al.’s terminology to our context, it was assumed that the adversary *knew* all of the claims in a given tree, and after the user revealed some claims, the adversary would attempt to convince an identity provider that some of the remaining unrevealed claims were in the tree. The security theorem given by Johnson, et al., stated that the adversary could not succeed, despite knowing all of the claims. For unforgeability, we need only a weaker version of this security that prevents an adversary from inserting new claims, and for privacy we must defend against adversaries who do not know all of the user’s claims and wish to learn something about the unrevealed claims. Second, combining trees did not apply in their context and was not considered. For this work, combining trees allows us to save the computation involved in verifying many signatures if the key used to sign the root node is trusted.

In the following definitions, we treat a credential system as a pair of algorithms (KeyGen, SignCred). KeyGen takes as input a security parameter and outputs a unique ID for the identity provider and a public/secret key pair to be used in generating and verifying credentials. SignCred takes as input an identity provider’s secret key and a set of claims, and outputs a credential on the set of claims. Each claim is either an arbitrary string s or a tuple (id, s, σ) , where id is

an identity provider's unique ID, s is an arbitrary string, and σ is a signature. The second type of claim corresponds to a request to combine a subtree with root s from the identity provider with unique ID id , and σ is the signature on s .

For the following definition we define two oracles, `CreateIDP` and `RetrieveCred`. `CreateIDP` runs `KeyGen` and returns the id and public key, and `RetrieveCred` takes as input a claim set and an identity provider's unique ID and runs `SignCred` with the corresponding secret key, and then returns the credential.

Definition 1. (Unforgeability) Let $(\text{KeyGen}, \text{SignCred})$ be a credential system. Then a probabilistic, poly-time adversary $\mathbf{A}^{\text{CreateIDP}, \text{RetrieveCred}}$ is said to forge a credential if it outputs a tuple $(pk^*, \sigma^*, \{c_i^*\}, \{n_i^*\})$, where pk^* is a public key output by `GenIDP`, σ^* is a valid signature under pk^* , $\{c_i^*\}$ is a set of claims not signed by their respective identity providers, and $\{n_i^*\}$ is a valid set of intermediate nodes for revealing the claims in a hash tree corresponding to σ^* .

Definition 2. (Privacy) Let $(\text{KeyGen}, \text{SignCred})$ be a credential system, let `Chal` be an oracle that when given two claim sets and a identity provider's ID, chooses one of the sets at random and generates a credential on that set, and let `Reveal` be an oracle that takes as input a claim set and reveals those claims (`Reveal` may only be called after `Chal` has been called). Then an adversary $\mathbf{A}^{\text{CreateIDP}, \text{KeyGen}, \text{Chal}, \text{Reveal}}$ is said to violate the privacy of the credential system if it guesses the set chosen by `Chal` with probability non-negligibly greater than $1/2$.

THEOREM 1. *If h is a collision resistant hash function and the underlying signature scheme is existentially unforgeable, then the scheme described above is unforgeable except with negligible probability.*

PROOF. Above we assumed that a tree is valid only if all of its claims are contained in subtrees or if none of its claims are in subtrees. We will deal with these two cases separately.

Case 1: \mathbf{A} outputs no subtree claims. Here σ^* must be a signature output by the underlying signature scheme under a key output by `CreateIDP`, because otherwise \mathbf{A} could be used to break the underlying signature scheme. Moreover, the nodes revealed by \mathbf{A} must be a subtree of the tree from that query. Otherwise \mathbf{A} has found a collision in the hash function: either a tree node's input was changed, or if an internal node was used as a leaf, then the appended bit is different, forcing a collision.

Case 2: \mathbf{A} outputs a tree containing subtrees. Here again the signature at the root provided by \mathbf{A} must have been output by `SignCred`. \mathbf{A} must reveal the intermediate nodes leading up to the root, and since every claim is contained in a subtree, \mathbf{A} must reveal a preimage of the trinary node (the root of the subtree). By the collision resistance of H , \mathbf{A} must reveal the same public key as part of the preimage of the trinary node. By the same reasoning, the revealed tree must be contained in the honestly generated tree, so the claims must have been signed before and the matching trinary node preimage guarantees that the claim is under the original identity provider. This completes the proof sketch. \square

We note that the random padding is not needed for unforgeability, but it is necessary for privacy, as discussed below.

Next we argue that the scheme does not leak information about unrevealed claims. Below we will need that the hash function is a pseudorandom function when the random padding is viewed as a key.

THEOREM 2. *If $H_k(x) = H(k, x)$ is a pseudorandom function (keyed with the padding k), then the scheme described above is private.*

PROOF. Consider an adversary \mathbf{A} attempting to violate the privacy of the credential system, as defined above. \mathbf{A} will obtain credentials on sets until it outputs two challenge sets, which define a symmetric difference of claims which it cannot ask to be revealed.

We define a sequence of "hybrid" versions of the game played by \mathbf{A} , where in each hybrid we hash one more claim in the symmetric difference with an independent random function instead of with H_k . Since H_k is a pseudorandom function that is only queried once for each randomly chosen k , \mathbf{A} cannot detect the difference between adjacent hybrids. Then \mathbf{A} cannot detect the difference between the extreme hybrids, where either all claims in the symmetric difference are hashed with H_k or all are independent random values. But in the latter case, the distribution seen by \mathbf{A} is the same no matter which of the two challenge claim sets is chosen, meaning that \mathbf{A} has advantage 0 in that hybrid. Thus \mathbf{A} has negligible advantage in distinguishing the two claim sets, completing the proof sketch. \square